## Recursive filtering in Halide

Dillon Sharlet, Google

## Description

- Recursive filter for a 1D signal

$$
\mathbf{y}_{\mathrm{n}}=(1-A) \mathbf{y}_{\mathrm{n}-1}+\mathrm{A} \mathbf{x}_{\mathrm{n}}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

## Description

- Recursive filter for a 1D signal

$$
y_{n}=(1-A) y_{n-1}+A x_{n}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function



## Description

- Recursive filter for a 1D signal

$$
y_{n}=(1-A) y_{n-1}+A x_{n}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function



## Description

- Recursive filter for a 1D signal

$$
y_{n}=(1-A) y_{n-1}+A x_{n}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function



## Description

- Recursive filter for a 1D signal

$$
y_{n}=(1-A) y_{n-1}+A x_{n}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function



## Description

- Recursive filter for a 1D signal

$$
y_{n}=(1-A) y_{n-1}+A x_{n}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function



## Description

- Recursive filter for a 1D signal

$$
y_{n}=(1-A) y_{n-1}+A x_{n}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function



## Description

- Recursive filter for a 1D signal

$$
y_{n}=(1-A) y_{n-1}+A x_{n}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function



## Description

- Recursive filter for a 1D signal

$$
y_{n}=(1-A) y_{n-1}+A x_{n}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function



## Description

- Recursive filter for a 1D signal

$$
y_{n}=(1-A) y_{n-1}+A x_{n}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function



## Description

- Recursive filter for a 1D signal

$$
y_{n}=(1-A) y_{n-1}+A x_{n}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function



## Description

- Recursive filter for a 1D signal

$$
y_{n}=(1-A) y_{n-1}+A x_{n}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function



## Description

- Recursive filter for a 1D signal

$$
\mathbf{y}_{\mathrm{n}}=(1-A) \mathbf{y}_{\mathrm{n}-1}+A \mathbf{x}_{\mathrm{n}}
$$

where $\mathbf{x}$ is input, $\mathbf{y}$ is output, $\mathbf{A}$ is the filter coefficient

- Example applied to a delta function

$$
\begin{aligned}
& \mathrm{x}=\square \square \square \square \square \square \square \square \\
& \mathrm{y}=\square \square \square \square \square \square \square \square \square \square \square \square \square ा
\end{aligned}
$$

## Description

- To apply this recursive filter to an image, apply it four times:
a. Up and down the columns
b. Right and left across the rows



## Reductions

- To implement this algorithm:
- Need to reference output at previous pixel to compute current output
- This cannot be done with a pure definition
- We can do this with update stages and RDoms
- RDom (Reduction Domain) provides a serial loop
- Can have dependencies between loop iterations


## Multi-stage Funcs

$$
\begin{aligned}
& f(x, y)=x+y ; \\
& f(x, 0)+=5 ;
\end{aligned}
$$

# Funcs can have multiple stages 

We call the additional ones
"update" stages

They run in sequence

$$
\begin{aligned}
& f(x, y)=x+y ; \\
& f(x, 0)+=5 ;
\end{aligned}
$$

They can use arbitrary index expressions on the left-hand-side

$$
\begin{array}{ll}
f(x, y)=x+y ; & \text { They } \\
f(x, 0)+=5 ; & \text { defin } \\
/ / f(x, 0)=f(x, 0)+5 ;
\end{array}
$$

$f(x, y)=x+y ;$
$f(x, 0)+=5$;
f.vectorize(x, 8);

They are scheduled independently
$f(x, y)=x+y ;$
$f(x, 0)+=5 ;$
f.vectorize(x, 8);
f.update(0)
.unroll(x, 2);

They are scheduled independently

$$
\begin{array}{ll}
f(x, y)=x+y ; & \text { for } y: \\
f(x, 0)+=5 ; & \text { for } x: \\
& f[x, y]=x+y
\end{array}
$$

$$
\begin{array}{ll}
f(x, y)=x+y ; & \text { for } y: \\
f(x, 0)+=5 ; & \text { for } x: \\
& f[x, y]=x+y \\
& \text { for } x: \\
& f[x, 0]=f[x, 0]+5 ;
\end{array}
$$

$f(x, y)=x+y ;$
RDom r(1, 10);
$f(x, 0)+=f(x, r)$;

An update stage can be a reduction over some domain "RDom"
$f(x, y)=x+y ;$ RDom $r(1,10)$; $f(x, 0)+=f(x, r)$;

This just throws an extra loop around the loop nest for that stage:
for $r$ from 1 to 10:
for $x$ :

$$
f[x, 0]=f[x, 0]+f[x, r] ;
$$

$f(x, y)=x+y ;$
RDom $r(1,10)$;

You can schedule RDom variables
$f(x, 0)+=f(x, r)$;
f.update(0)
.unroll(r);
$f(x, y)=x+y ;$
RDom $r(1,10)$;

You can schedule RDom variables
$f(x, 0)+=f(x, r)$;
f.update(0)
.reorder (r, x);
$f(x, y)=x+y ;$ RDom $r(1,10)$; $f(x, 0)+=f(x, r)$; f.update(0)
.parallel(r);

## ERROR: Potential

Race Condition

But only when we can prove there's no race condition or change in meaning.

Halide's promise:
Scheduling never changes the results!

## Generators

- Two ways to call Halide code
- JIT: Halide pipelines executed in the same process they are defined in
- AOT: Halide pipelines compiled to object files (.o, . obj) and linked into/called from another program via C ABI (i.e. extern "C")


## Generators

- Generators are C++ programs that, when run, produce objects (.o, .obj) and C headers (.h) containing compiled pipelines
- Applications \#include generated header files declaring the functions, link to generated objects
- Pipeline functions are declared with arguments corresponding to Param objects, including ImageParams in buffer t objects.
- Holds pointer, element size and strides of each dimension of an image
- Halide never assumes ownership of the memory a buffer_t points to


## Using Generators with Matlab

- Generators can also be used within Matlab (or Octave) via the mex library interface
- Halide pipeline compiled with mat lab target feature defines a suitable mexFunction wrapper
- Validates and converts mxArray to buffer_t (or scalar params)
- mex_halide Matlab function performs all the required steps to build a mex library from a source file containing a generator

Code!

## Scheduling for locality

- So far, we've talked about some scheduling operators
- vectorize, unroll, etc.
- We've also briefly discussed compute_at
- To significantly improve performance, we need to use compute_at to improve locality


## compute_root

$$
\begin{aligned}
& f(x, y)=x+y ; \\
& g(x, y)=2 * f(x, y) ;
\end{aligned}
$$

Here is a simple two stage pipeline

## compute_root

$f(x, y)=x+y ;$
$g(x, y)=2 * f(x, y)$; f.compute_root();
g.compute_root();

This means compute all of $f$, followed by all of g

Poor locality!

## compute_root

$f(x, y)=x+y ; \quad$ for $f . y:$
$g(x, y)=2 * f(x, y) ; \quad$ for $f . x$ : f.compute_root();
g.compute_root(); for g.y:
for $g . x$ :

$$
g[g \cdot x, g \cdot y]=2 * f[g \cdot x, g \cdot y]
$$

## compute_at

$f(x, y)=x+y ;$
$g(x, y)=2 * f(x, y)$; f.compute_at(g, y);
g.compute_root();
"Compute f at each iteration of $y$ when computing g"

All stages of a Func share the same compute_at location

## compute_at

$f(x, y)=x+y ; \quad$ for $g . y:$
$g(x, y)=2 * f(x, y)$; for $g . x$ :
f.compute_at(g, y); g[g.x,g.y] = 2*f[g.x,g.y] g.compute_root();

## compute_at

$f(x, y)=x+y ; \quad$ for $g \cdot y$ :
$g(x, y)=2 * f(x, y) ; \quad$ for $f . x$ :
f.compute_at (g, y);
g.compute_root();

$$
f[f \cdot x, g \cdot y]=f \cdot x+g \cdot y
$$

for $g . x$ :

$$
g[g \cdot x, g \cdot y]=2 * f[g \cdot x, g \cdot y]
$$

## compute_at

$$
\begin{array}{ll}
f(x, y)=x+y ; & \text { for } g \cdot y: \\
g(x, y)=2 * f(x, y) ; & \text { for } g \cdot x:
\end{array}
$$

f.compute_at (g, x), g[g.x,g.y] = 2*f[g.x,g.y] g.compute_root();

## compute_at

$f(x, y)=x+y ; \quad$ for $g \cdot y:$
$g(x, y)=2 * f(x, y)$; for $g \cdot x$ :
f.compute_at (g, x);
g.compute_root();

$$
\begin{aligned}
& f[g \cdot x, g \cdot y]=g \cdot x+g \cdot y \\
& g[g \cdot x, g \cdot y]=2 * f[g \cdot x, g \cdot y]
\end{aligned}
$$

## IIR blur compute_root visualization



## IIR blur locality schedule visualization

Transpose
Blur y


## IIR blur locality schedule visualization

Transpose
Blur y


## IIR blur locality schedule visualization

Transpose
Blur y


## IIR blur locality schedule visualization

Transpose
Blur y


## Code!

https://github.com/halide/CVPR2015/tree/master/RecursiveFilter

