

# An introduction to Halide

Jonathan Ragan-Kelley (Stanford)

Andrew Adams (Google)

Dillon Sharlet (Google)

# Today's agenda

Now: **the big ideas in Halide**

Later: **writing & optimizing real code**

Hello world (brightness)

Gaussian blur - *3x OpenCV*

Simple enhancement pipeline - *6x OpenCV*

MATLAB integration

IIR filter

CNN layers

GPU scheduling

*break*

*break*

Finally: **real-time HOG on a phone**

# **We are surrounded by computational cameras**

**Enormous opportunity,  
demands extreme optimization  
parallelism & locality limit  
performance and energy**


# We are surrounded by computational cameras

Enormous opportunity,  
demands extreme optimization  
parallelism & locality limit  
performance and energy

**Camera:** 8 Mpixels  
(96MB/frame as *float*)

**CPUs:** 15 GFLOP/sec

**GPU:** 115 GFLOP/sec





# We are surrounded by computational cameras


Enormous opportunity,  
demands extreme optimization  
parallelism & locality limit  
performance and energy

**Camera:** 8 Mpixels  
(96MB/frame as *float*)

**CPUs:** 15 GFLOP/sec

**GPU:** 115 GFLOP/sec

***Required  
arithmetic  
intensity*** > 40:1



# Today's methodology

**C++ w/multithreading, SIMD**

**CUDA/OpenCL**

**OpenGL/RenderScript**

# Today's methodology

**C++ w/multithreading, SIMD**

**CUDA/OpenCL**

**OpenGL/RenderScript**

Optimization requires manually  
**transforming program & data structure**  
for locality and parallelism.

# Today's methodology

**C++ w/multithreading, SIMD**

**CUDA/OpenCL**

**OpenGL/RenderScript**

Optimization requires manually  
**transforming program & data structure**  
for locality and parallelism.

*libraries don't solve this:*

**BLAS, IPP, MKL, OpenCV**

optimized kernels compose into  
inefficient pipelines (no fusion)

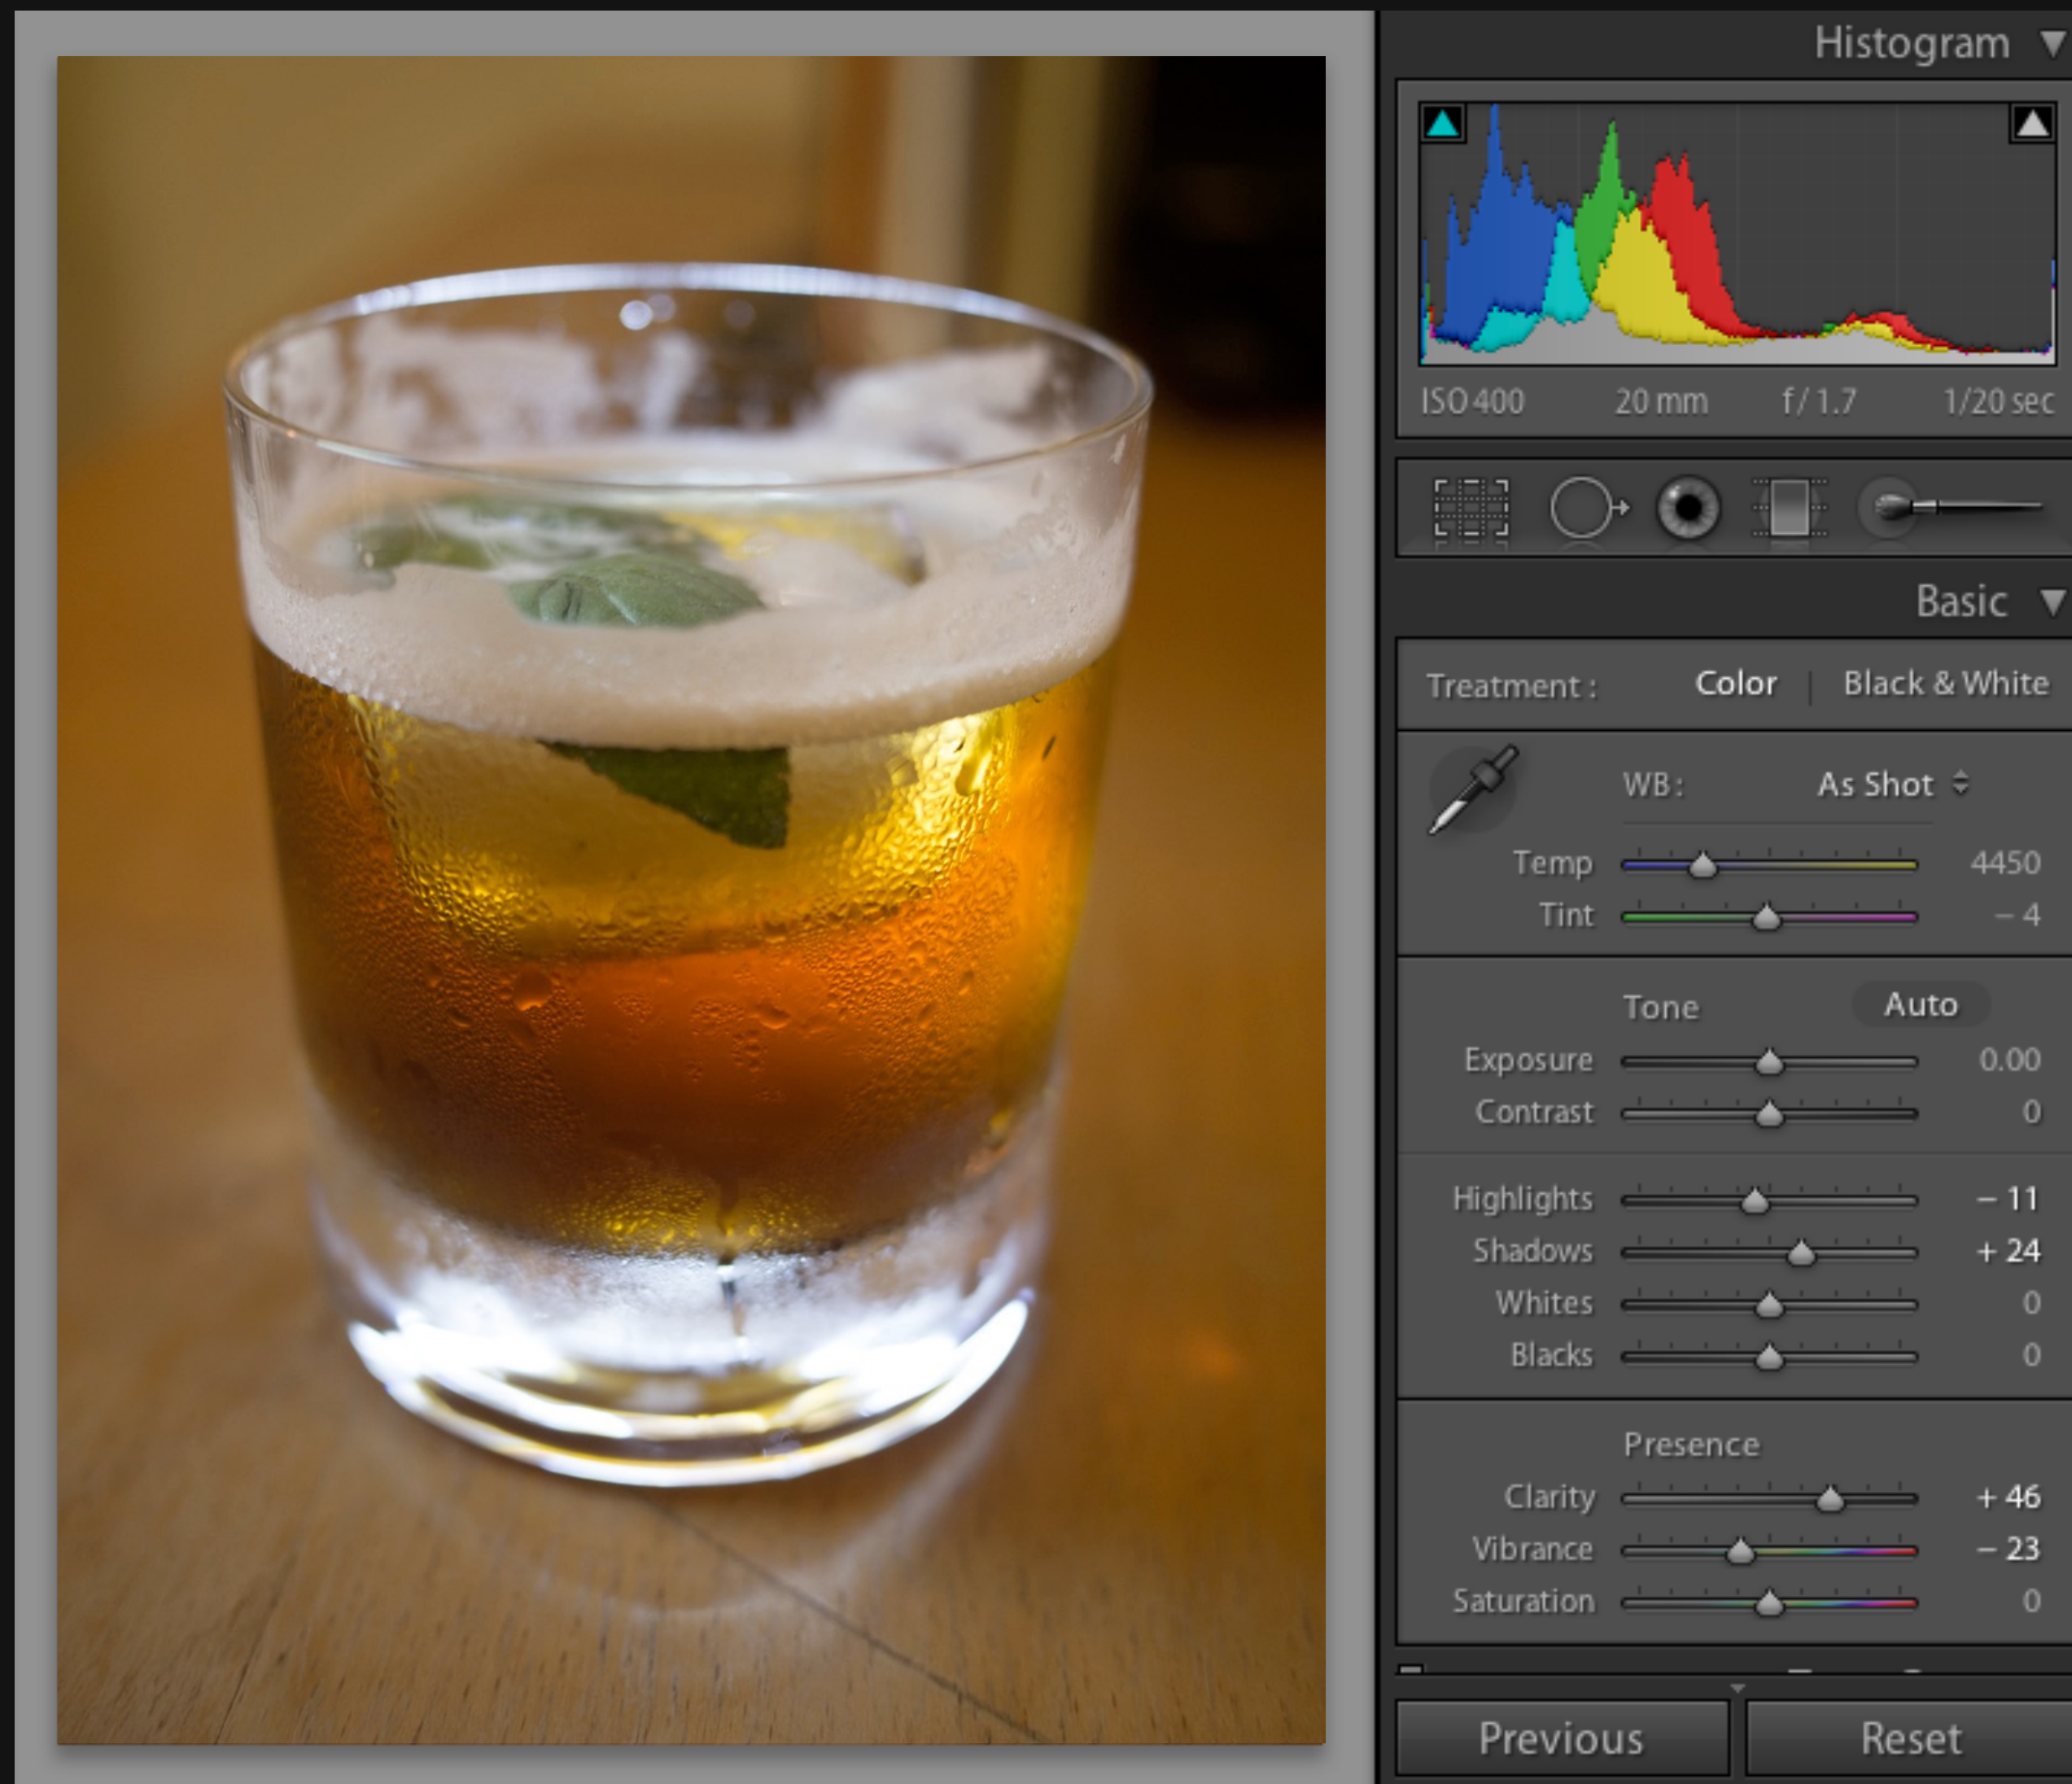




# Local Laplacian Filters

in Adobe Photoshop Camera Raw / Lightroom

**1500 lines of expert-  
optimized C++  
multi-threaded, SSE  
3 months of work  
10x faster than reference C**








# Local Laplacian Filters

in Adobe Photoshop Camera Raw / Lightroom

**1500 lines of expert-optimized C++**  
multi-threaded, SSE  
**3 months of work**  
**10x faster than reference C**





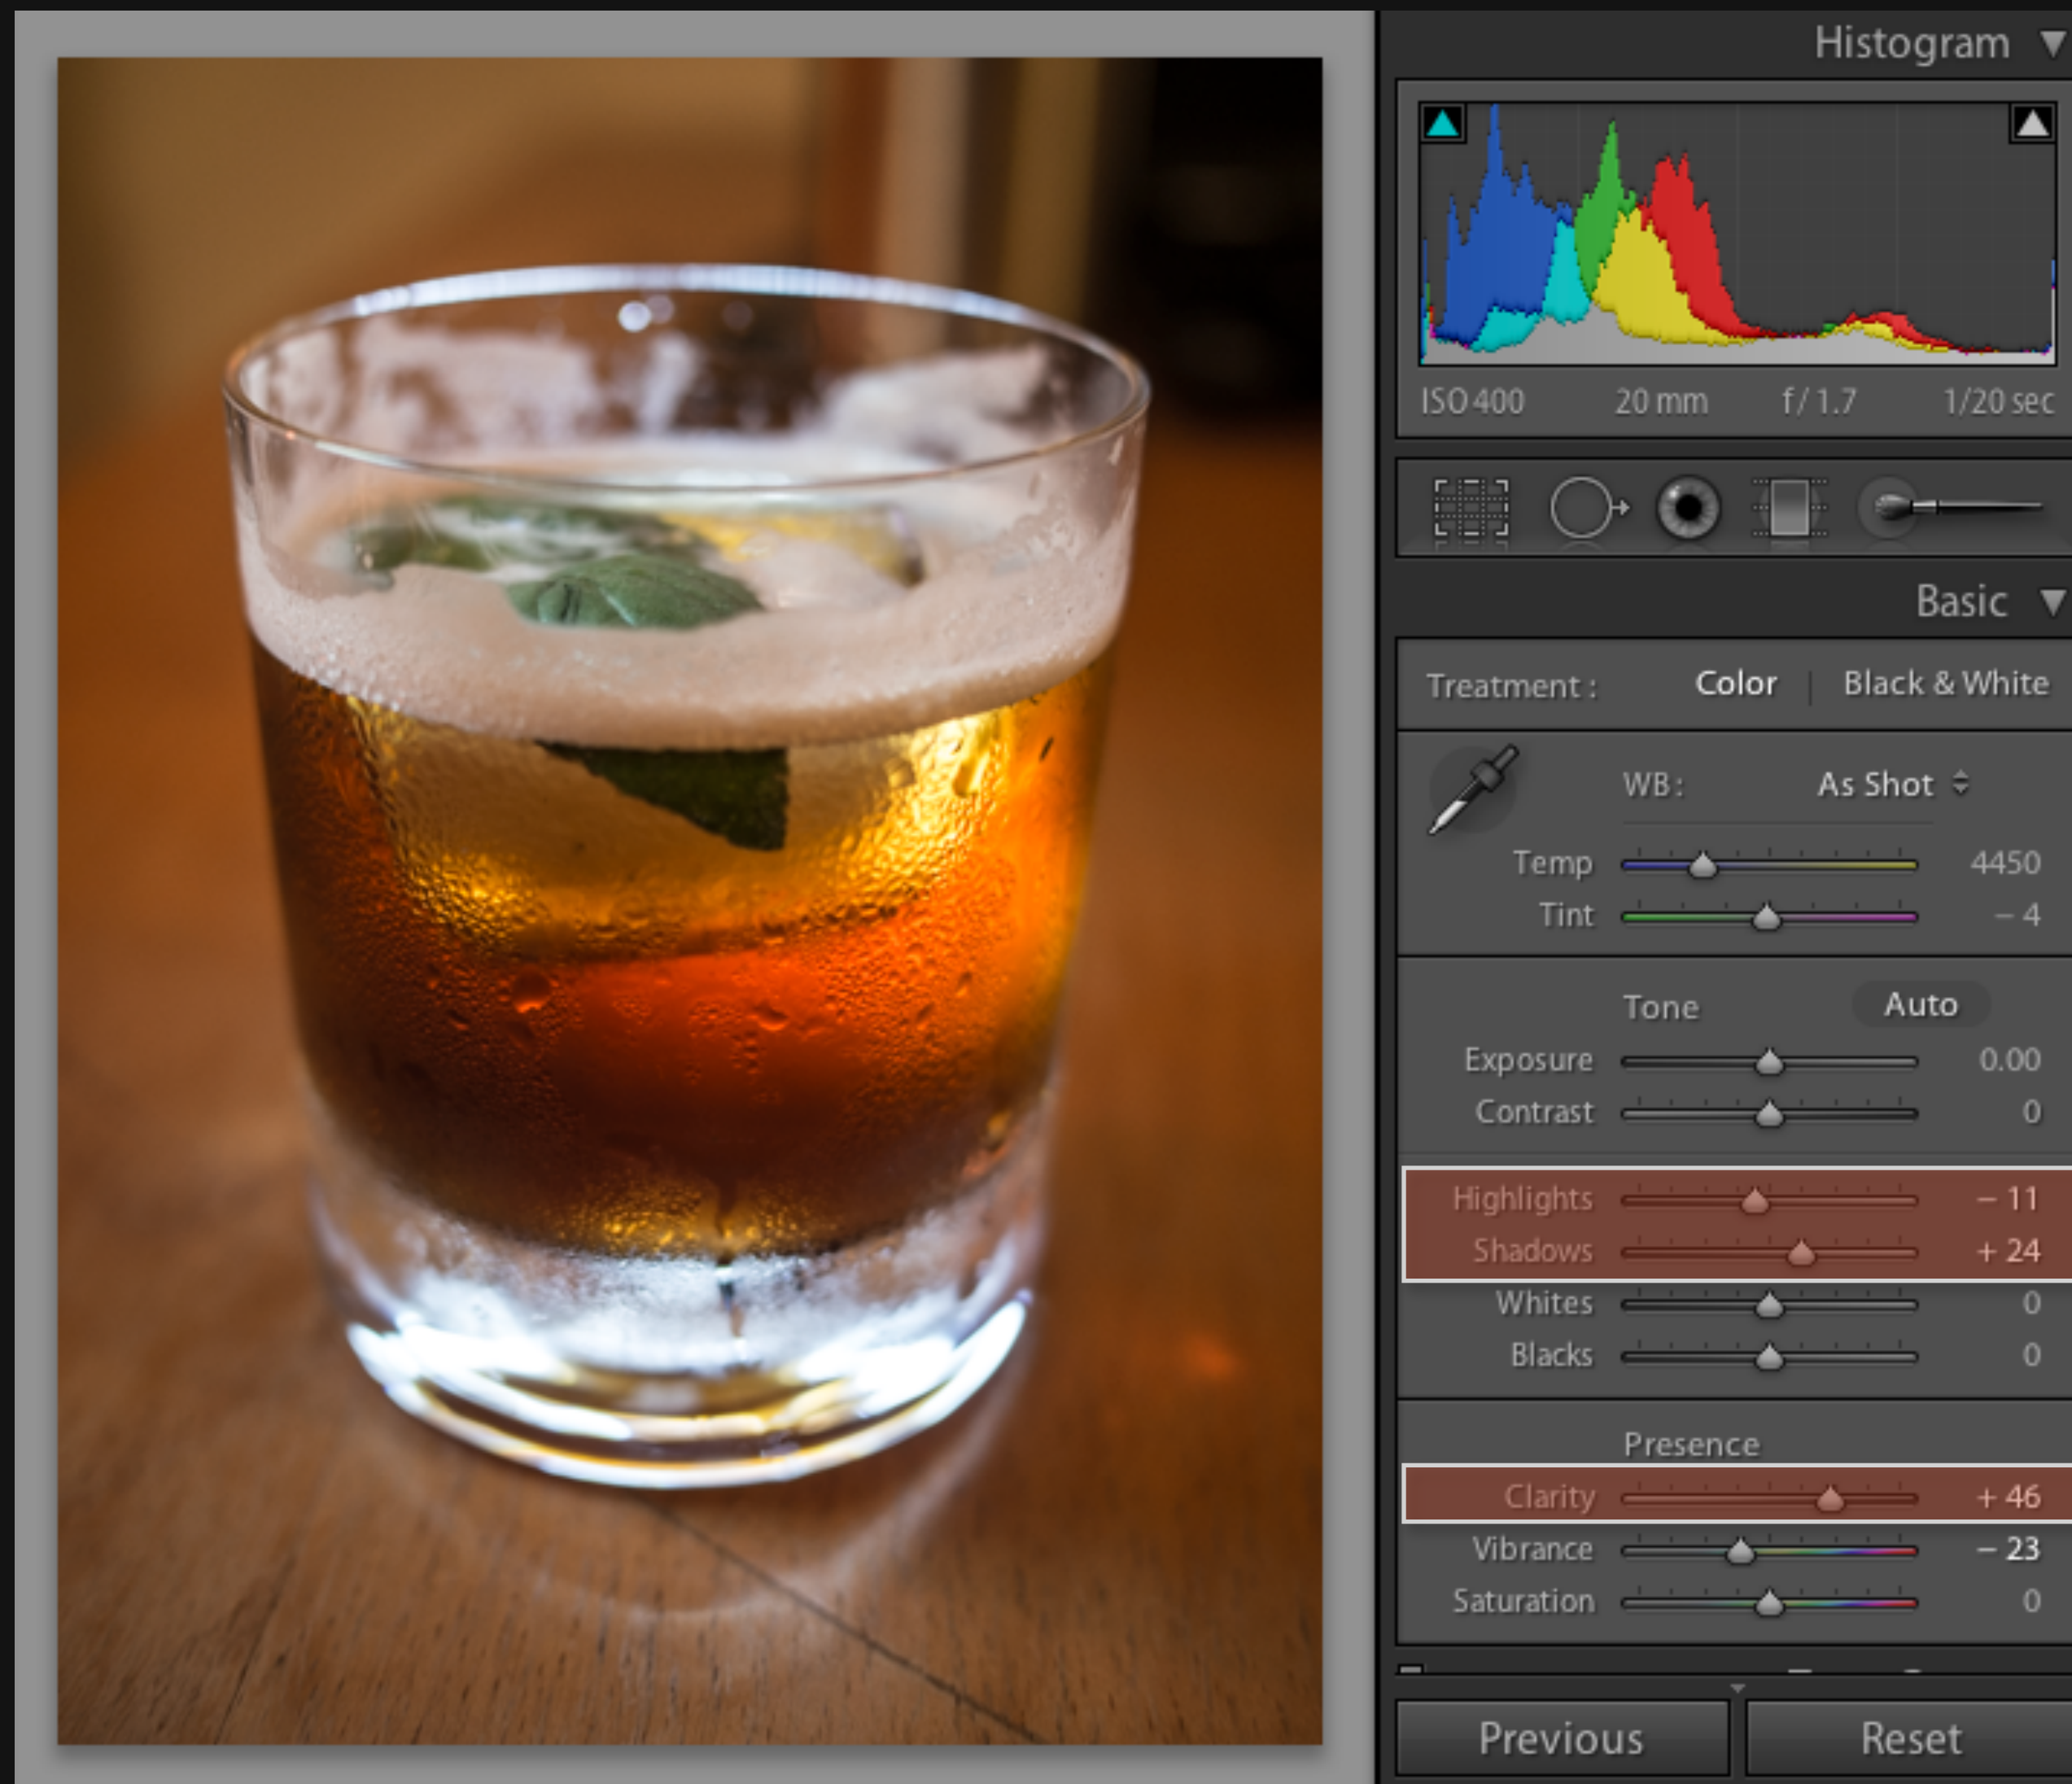


# Local Laplacian Filters

## in Adobe Photoshop Camera Raw / Lightroom

**1500 lines of expert-optimized C++**  
multi-threaded, SSE

***3 months of work***  
***10x faster than reference C***







# Local Laplacian Filters

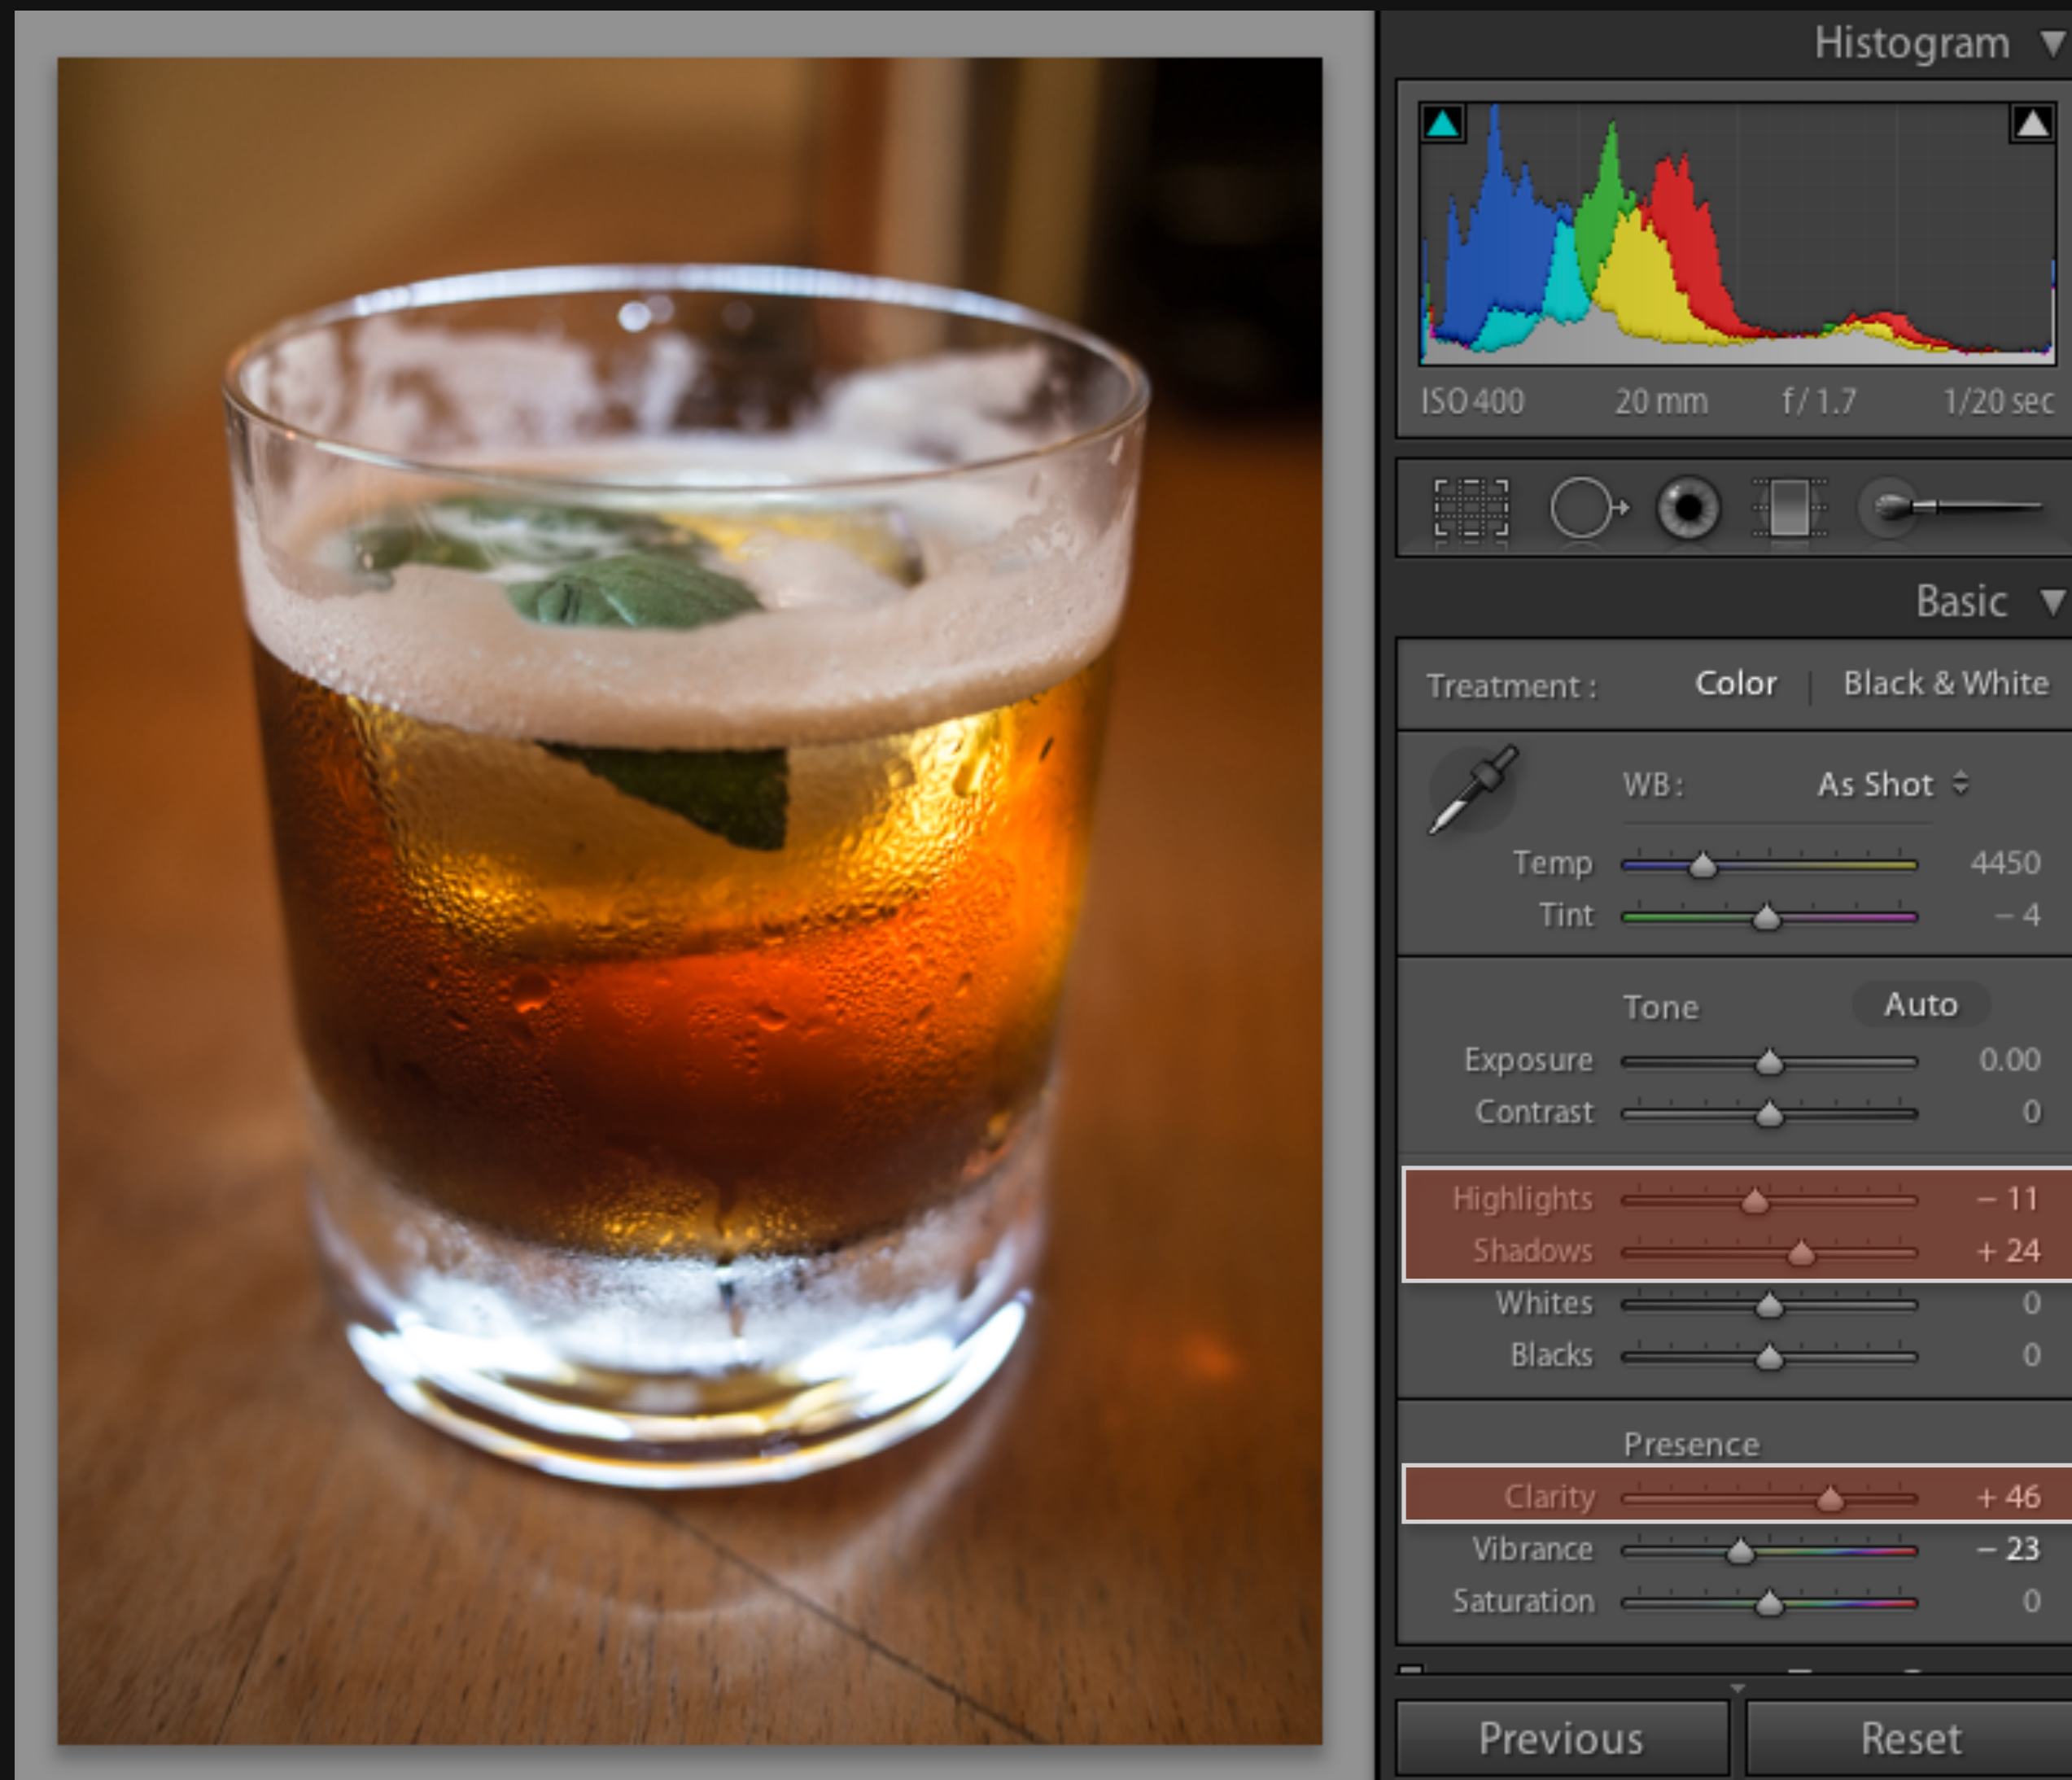
in Adobe Photoshop Camera Raw / Lightroom

1500 lines of expert-optimized C++ multi-threaded, SSE






*3 months of work*

*10x faster than reference C*



Just writing in C isn't nearly enough!







**Imaging is  
*everywhere***





# A simple example: 3x3 blur

```
void box_filter_3x3(const Image &in, Image &blury) {
    Image blurx(in.width(), in.height()); // allocate blurx array

    for (int y = 0; y < in.height(); y++)
        for (int x = 0; x < in.width(); x++)
            blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;

    for (int y = 0; y < in.height(); y++)
        for (int x = 0; x < in.width(); x++)
            blury(x, y) = (blurx(x, y-1) + blurx(x, y) + blurx(x, y+1))/3;
}
```

# Hand-optimized C++

9.9 → 0.9 ms/megapixel

```
void box_filter_3x3(const Image &in, Image &blury) {
    __m128i one_third = _mm_set1_epi16(21846);
    #pragma omp parallel for
    for (int yTile = 0; yTile < in.height(); yTile += 32) {
        __m128i a, b, c, sum, avg;
        __m128i blurx[(256/8)*(32+2)]; // allocate tile blurx array
        for (int xTile = 0; xTile < in.width(); xTile += 256) {
            __m128i *blurxPtr = blurx;
            for (int y = -1; y < 32+1; y++) {
                const uint16_t *inPtr = &(in[yTile+y][xTile]);
                for (int x = 0; x < 256; x += 8) {
                    a = _mm_loadu_si128((__m128i*)(inPtr-1));
                    b = _mm_loadu_si128((__m128i*)(inPtr+1));
                    c = _mm_load_si128((__m128i*)(inPtr));
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                    avg = _mm_mulhi_epi16(sum, one_third);
                    _mm_store_si128(blurxPtr++, avg);
                    inPtr += 8;
                }
            }
            blurxPtr = blurx;
            for (int y = 0; y < 32; y++) {
                __m128i *outPtr = (__m128i *)(&(blury[yTile+y][xTile]));
                for (int x = 0; x < 256; x += 8) {
                    a = _mm_load_si128(blurxPtr+(2*256)/8);
                    b = _mm_load_si128(blurxPtr+256/8);
                    c = _mm_load_si128(blurxPtr++);
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                    avg = _mm_mulhi_epi16(sum, one_third);
                    _mm_store_si128(outPtr++, avg);
                }
            }
        }
    }
}
```

**11x faster**  
(quad core x86)

Tiled, fused

Vectorized

Multithreaded

Redundant  
computation

*Near roof-line  
optimum*

# Halide's answer: *decouple* algorithm from schedule

**Algorithm:** *what* is computed

**Schedule:** *where* and *when* it's computed

**Easy for programmers to build pipelines**

simplifies algorithm code

improves modularity

**Easy for programmers to specify & explore optimizations**

fusion, tiling, parallelism, vectorization

can't break the algorithm

**Easy for the compiler to generate fast code**

**The algorithm** defines pipelines as pure functions

Pipeline stages are *functions* from coordinates to values

Execution order and storage are unspecified

# The algorithm defines pipelines as pure functions

Pipeline stages are *functions* from coordinates to values

Execution order and storage are unspecified

## 3x3 blur as a Halide *algorithm*:

```
Var x, y; Func blurx, blury;
```

```
blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;
```

```
blury(x, y) = (blurx(x, y-1) + blurx(x, y) + blurx(x, y+1))/3;
```

# **Domain scope of the programming model**

All computation is over **regular grids**.

Only **feed-forward pipelines**

Recursive/reduction computations are a (partial) escape hatch.

**Recursion must have bounded depth.**

# **Domain scope** of the programming model

All computation is over **regular grids**.

**not**  
**Turing**  
**complete**

Only **feed-forward pipelines**

Recursive/reduction computations are a (partial) escape hatch.

**Recursion** must have **bounded depth**.



# **Halide's answer: *decouple* algorithm from schedule**

**Algorithm: *what*** is computed

**Schedule: *where* and *when*** it's computed

**Easy for programmers to build pipelines**

simplifies algorithm code

improves modularity

**Easy for programmers to specify & explore optimizations**

fusion, tiling, parallelism, vectorization

can't break the algorithm

**Easy for the compiler to generate fast code**

# Halide's answer: *decouple* algorithm from schedule

Algorithm: *what* is computed

Schedule: *where* and *when* it's computed

**Easy for programmers to build pipelines**

simplifies algorithm code

improves modularity

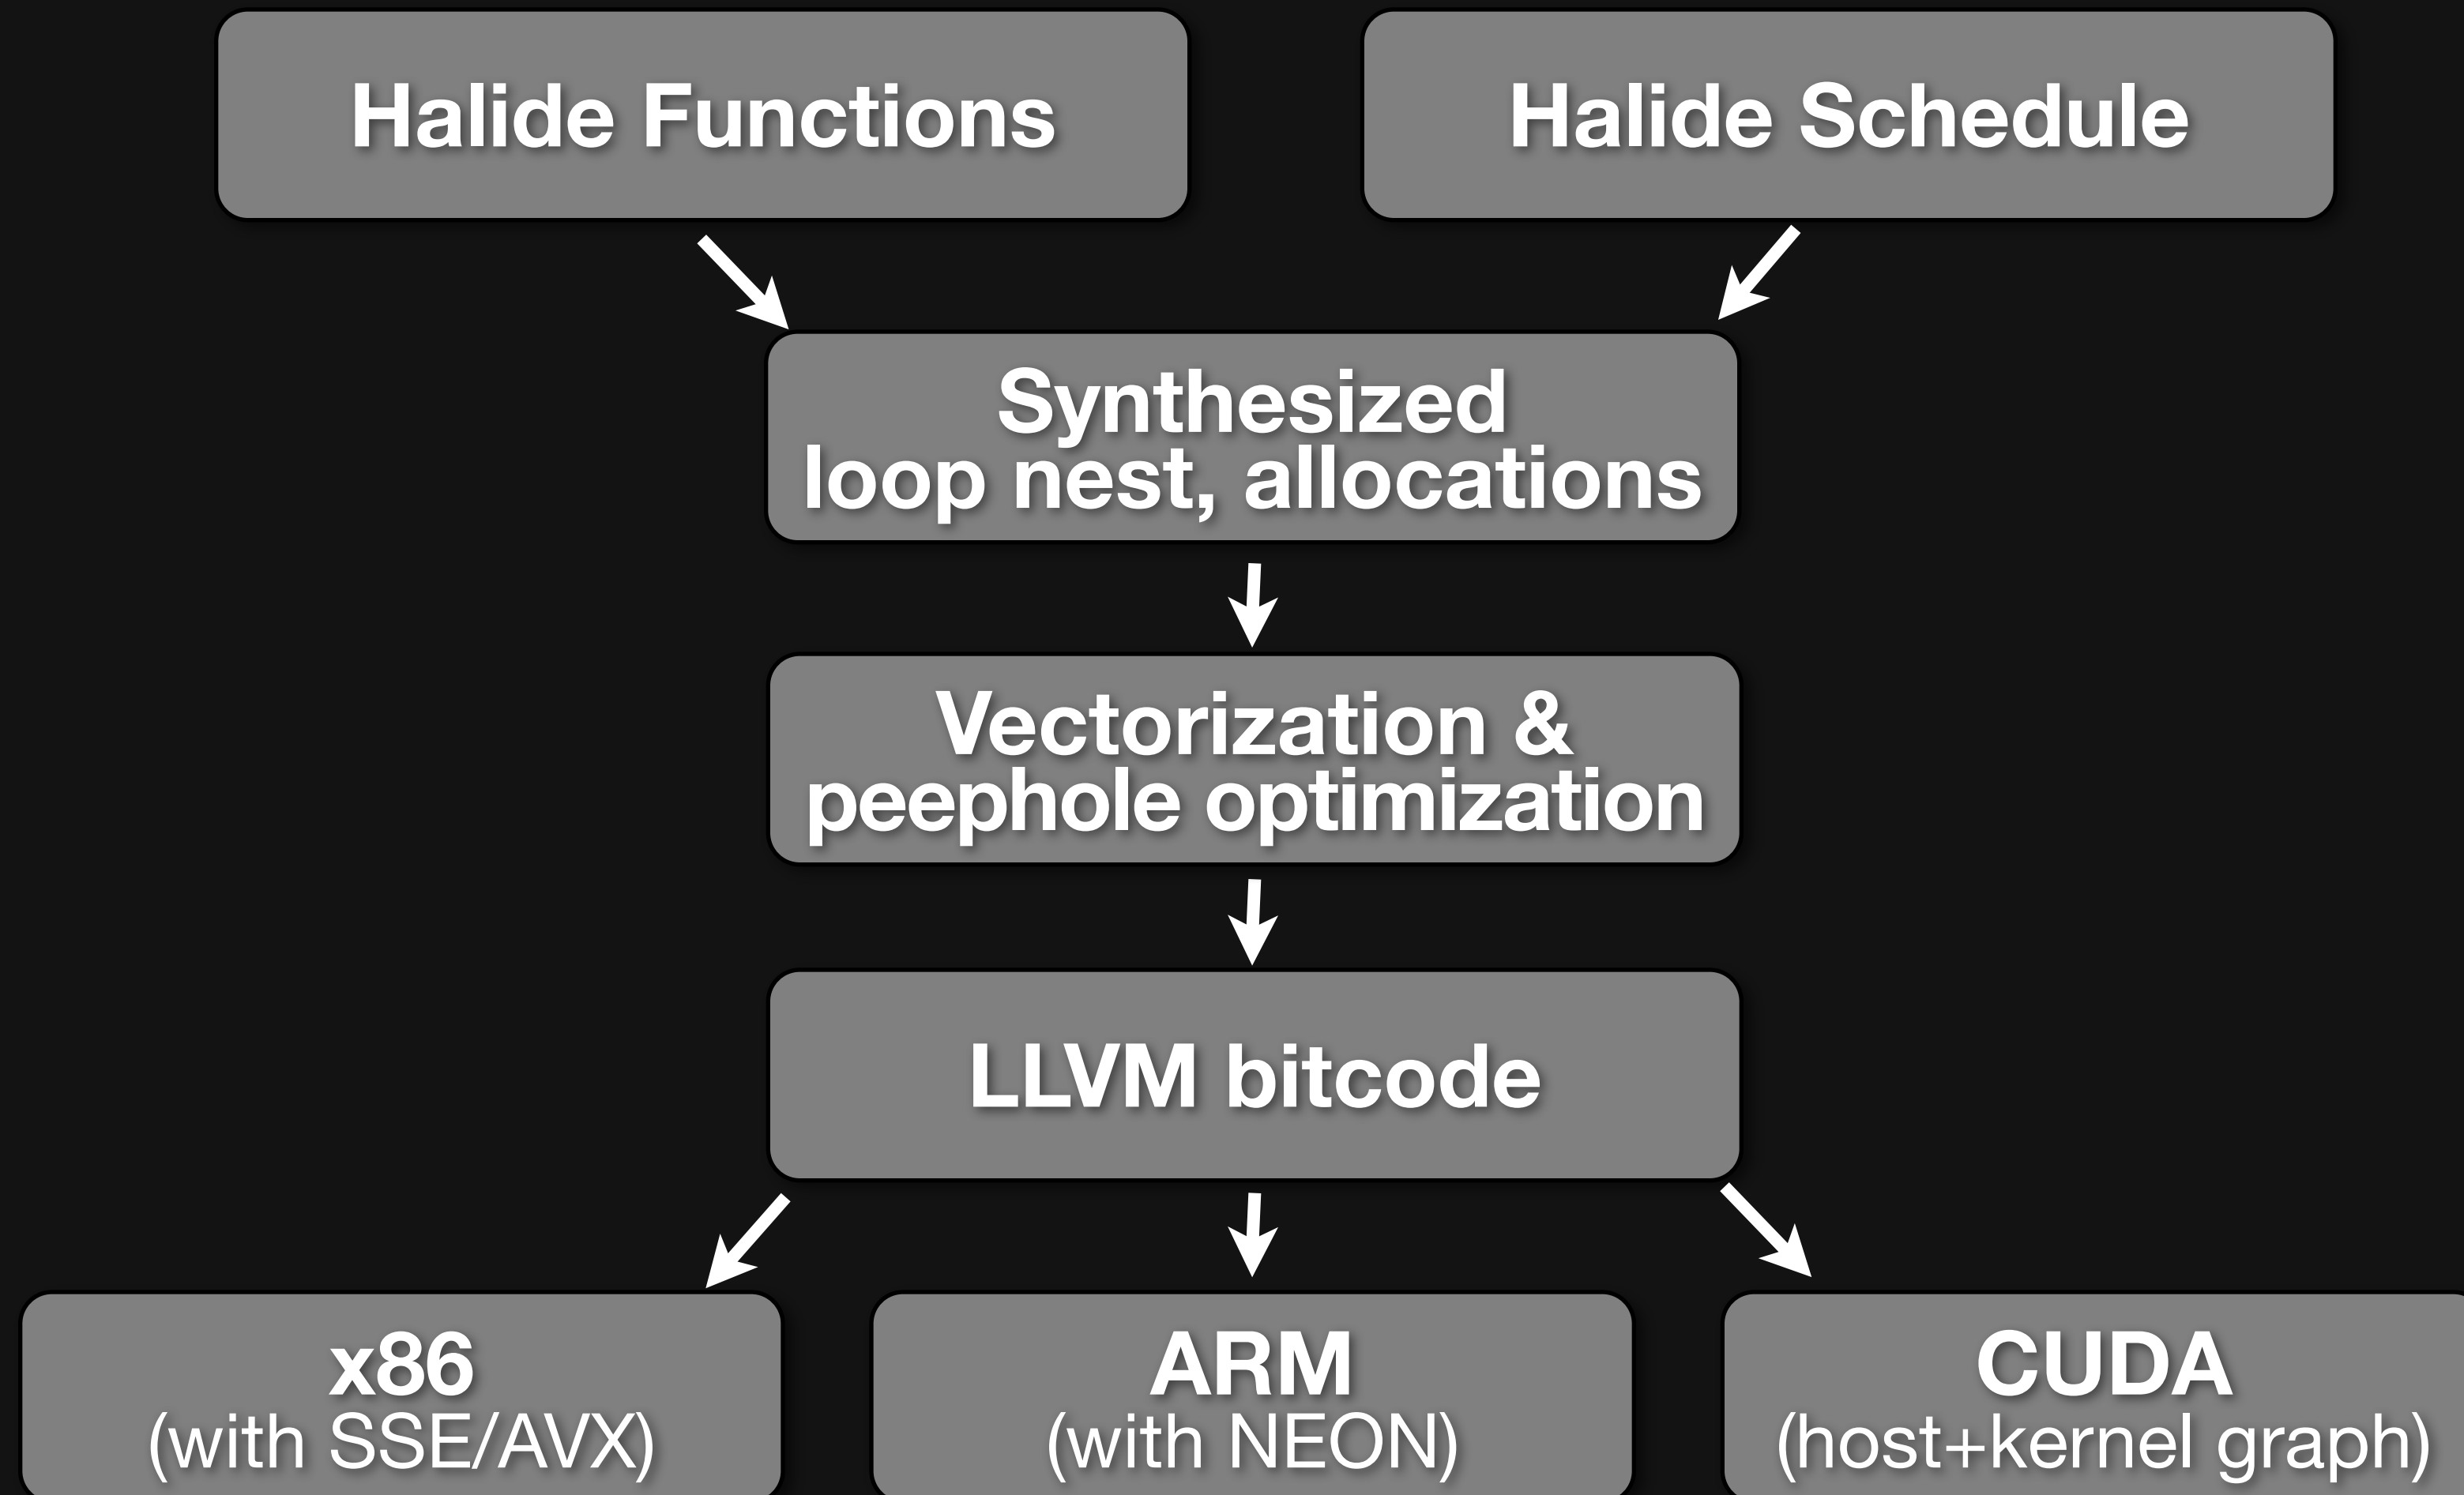
**Easy for programmers to specify & explore optimizations**

fusion, tiling, parallelism, vectorization

can't break the algorithm

**Easy for the compiler to generate fast code**

# The Halide Compiler



# The Halide Compiler

Halide Functions

Halide Schedule

Synthesized  
loop nest, allocations

Vectorization &  
peephole optimization

LLVM bitcode

NaCl, PNaCl  
(in-browser)

x86  
(with SSE/AVX)

ARM  
(with NEON)


CUDA, OpenCL, GL ES  
(host+kernel graph)

C  
(source)



# Local Laplacian Filters

prototype for Adobe Photoshop Camera Raw / Lightroom



Histogram ▼

ISO 400 20 mm f/1.7 1/20 sec

Basic ▼

Treatment: Color | Black & White

WB: As Shot ▼

Temp 4450

Tint -4

Tone Auto

Exposure 0.00

Contrast 0

Highlights -11

Shadows +24

Whites 0

Blacks 0

Presence

Clarity +46

Vibrance -23

Saturation 0

Previous Reset



# Local Laplacian Filters

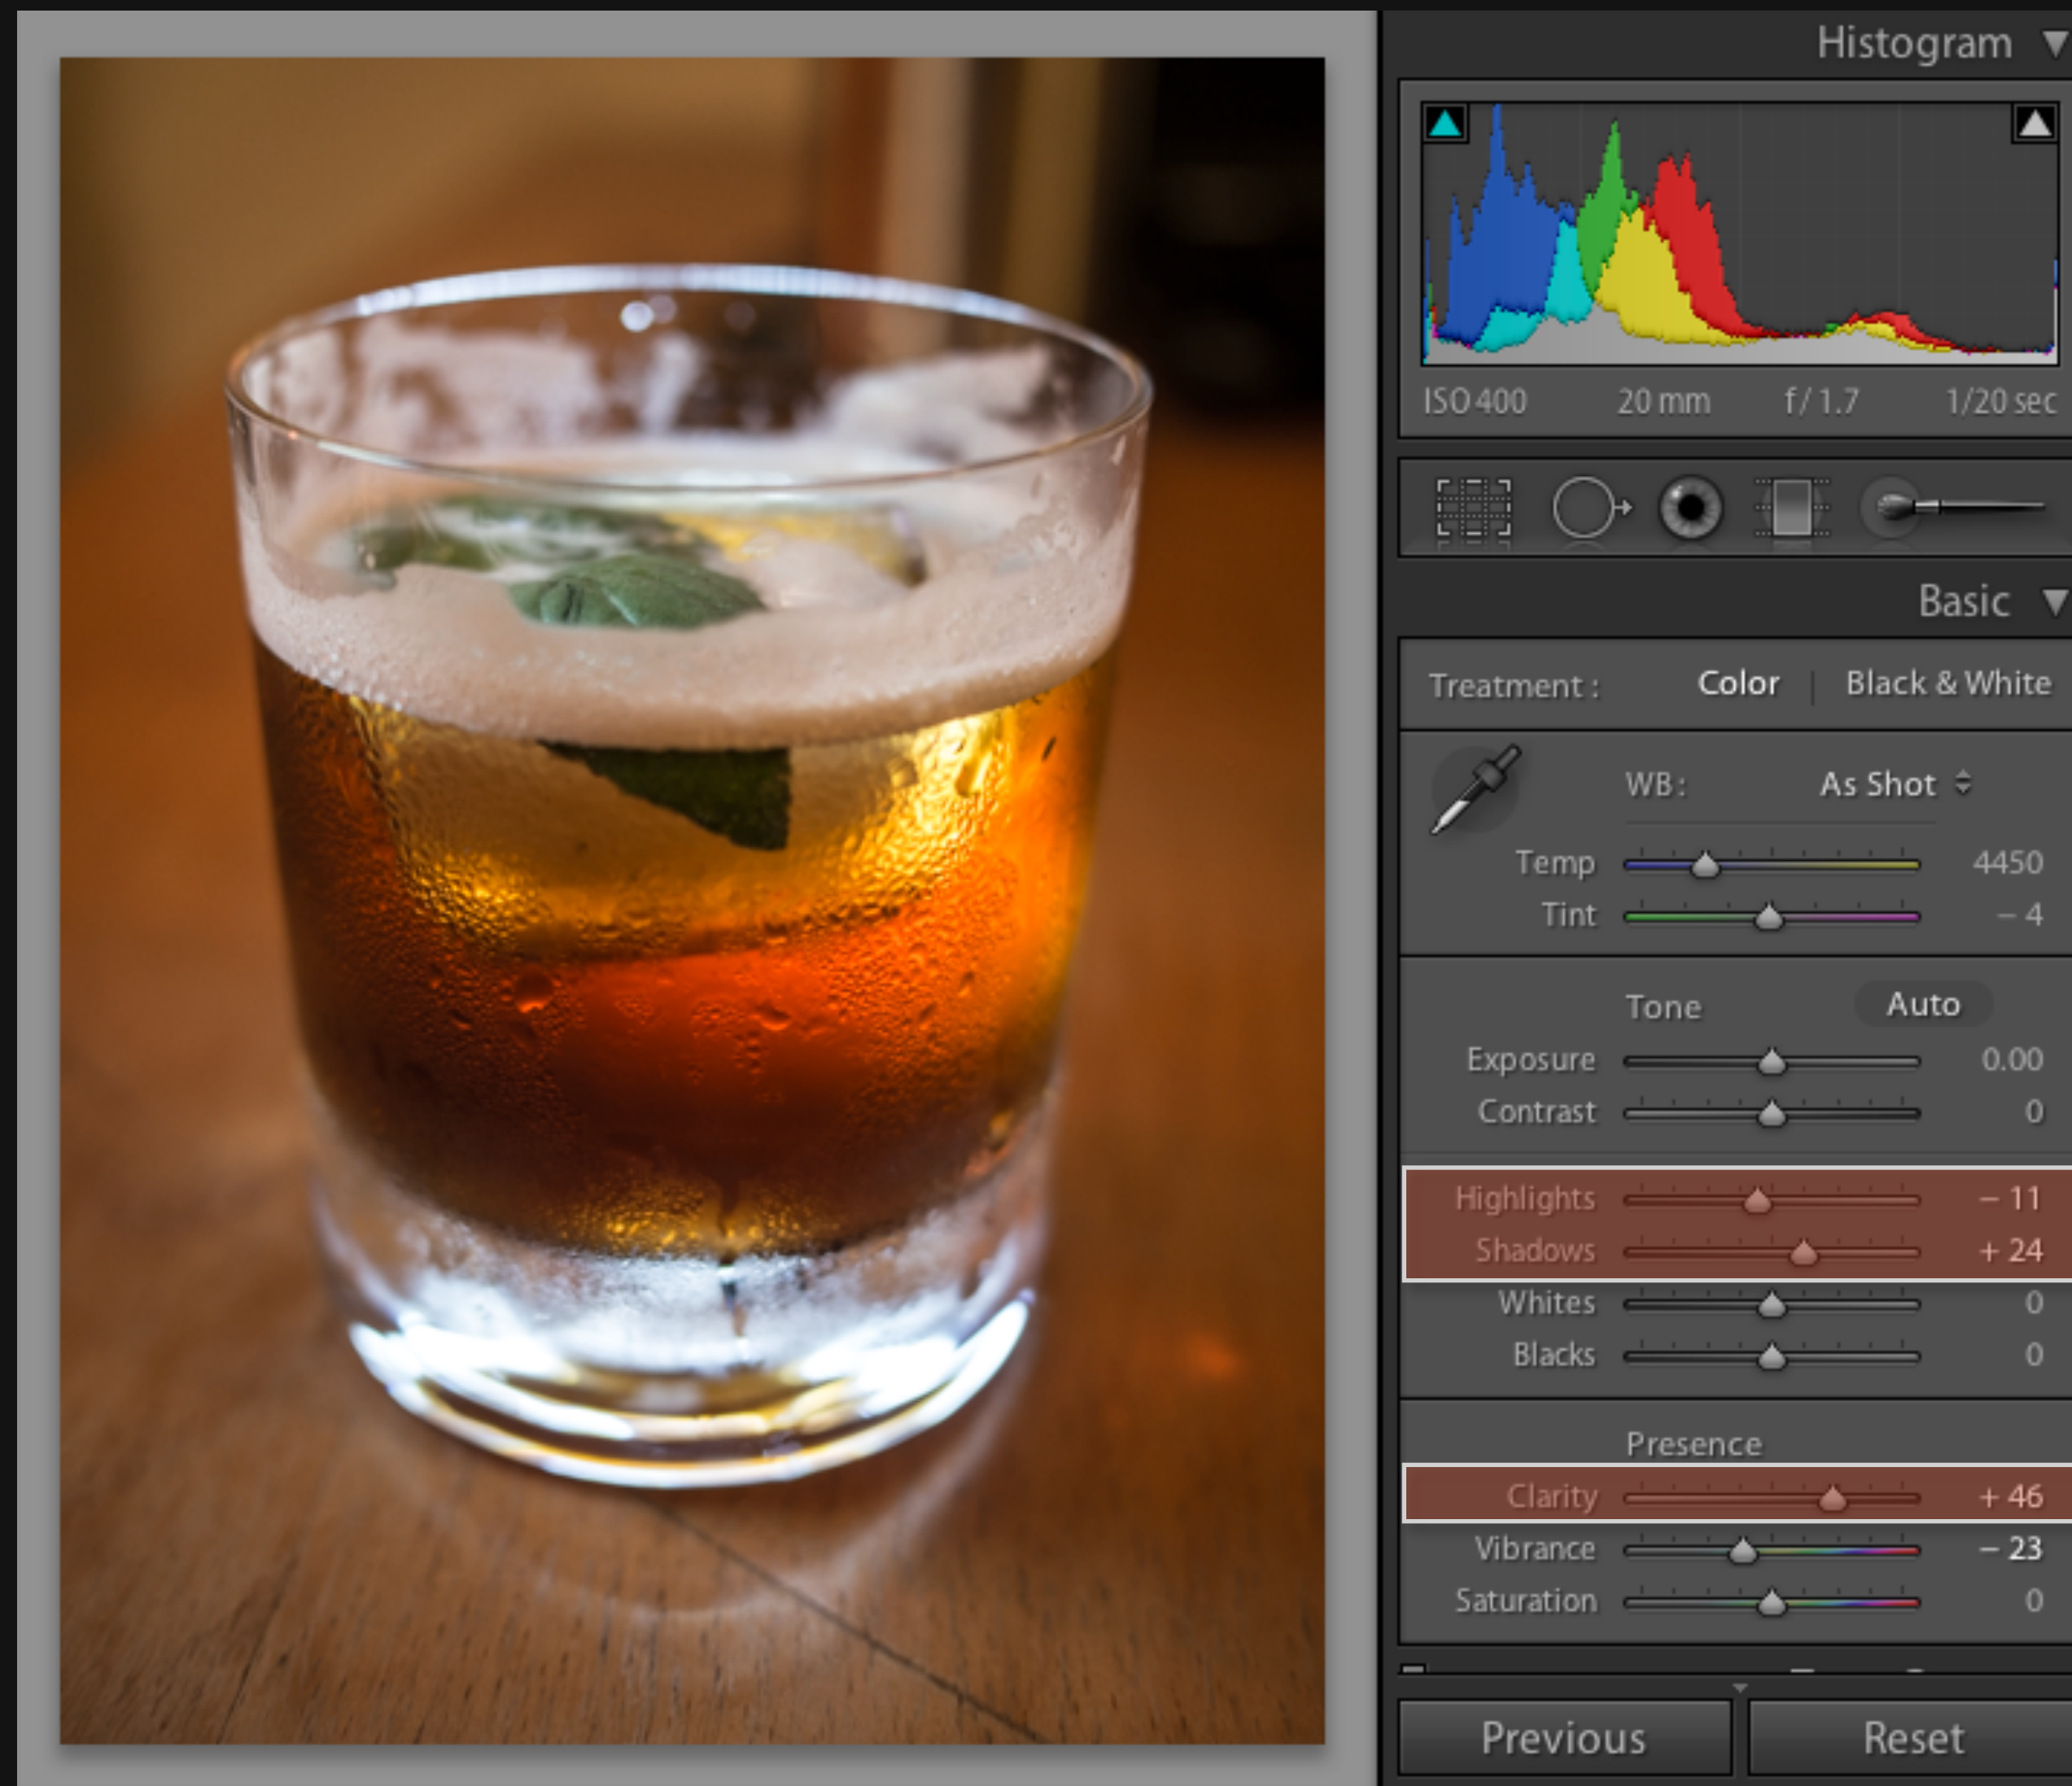
prototype for Adobe Photoshop Camera Raw / Lightroom

Reference: 300 lines C++

Adobe: 1500 lines

*3 months of work*

*10x faster (vs. reference)*





# Local Laplacian Filters

prototype for Adobe Photoshop Camera Raw / Lightroom

Reference: 300 lines C++

Adobe: 1500 lines

*3 months of work*


*10x faster (vs. reference)*

Halide: 60 lines

*1 intern-day*

*20x faster (vs. reference)*

*2x faster (vs. Adobe)*





# Local Laplacian Filters

prototype for Adobe Photoshop Camera Raw / Lightroom

**Reference: 300 lines C++**

**Adobe: 1500 lines**

*3 months of work*

*10x faster (vs. reference)*

**Halide: 60 lines**

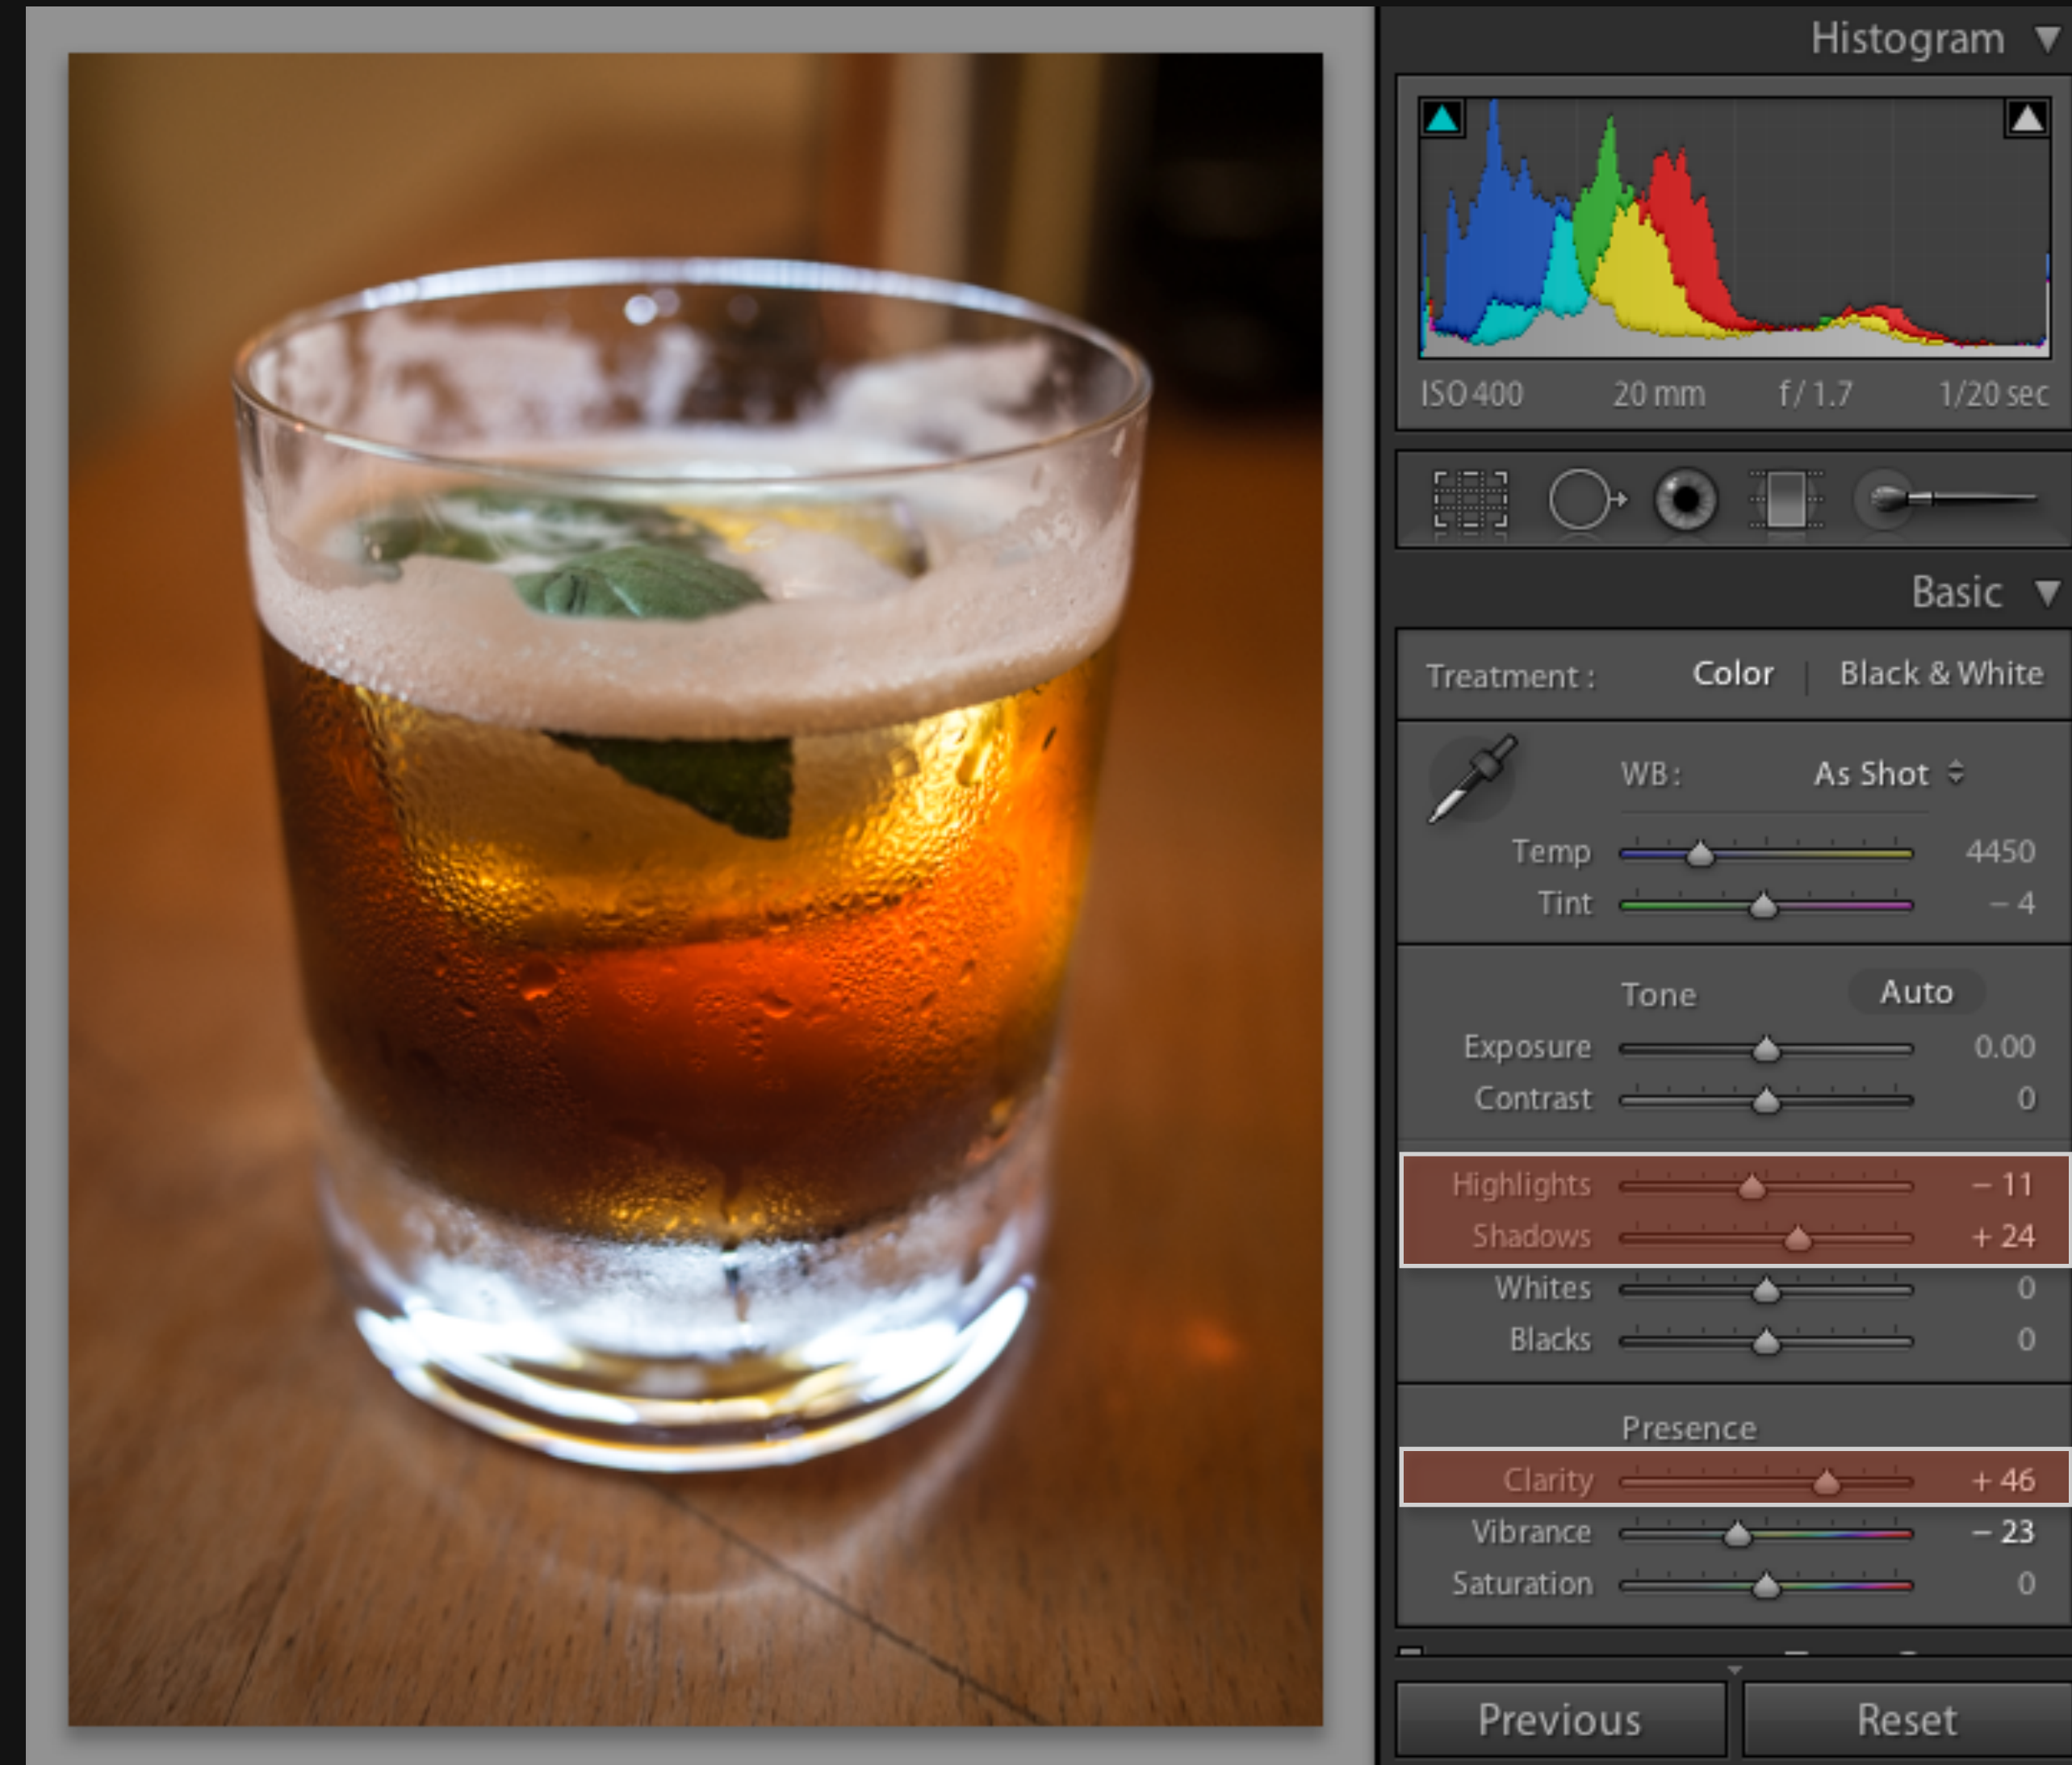
*1 intern-day*

**20x faster (vs. reference)**

**2x faster (vs. Adobe)**

**GPU: 90x faster (vs. reference)**

**9x faster (vs. Adobe)**





# Local Laplacian Filters

prototype for Adobe Photoshop Camera Raw / Lightroom

**Reference: 300 lines C++**

**Adobe: 1500 lines**

*3 months of work*

*10x faster (vs. reference)*

**Halide: 60 lines**

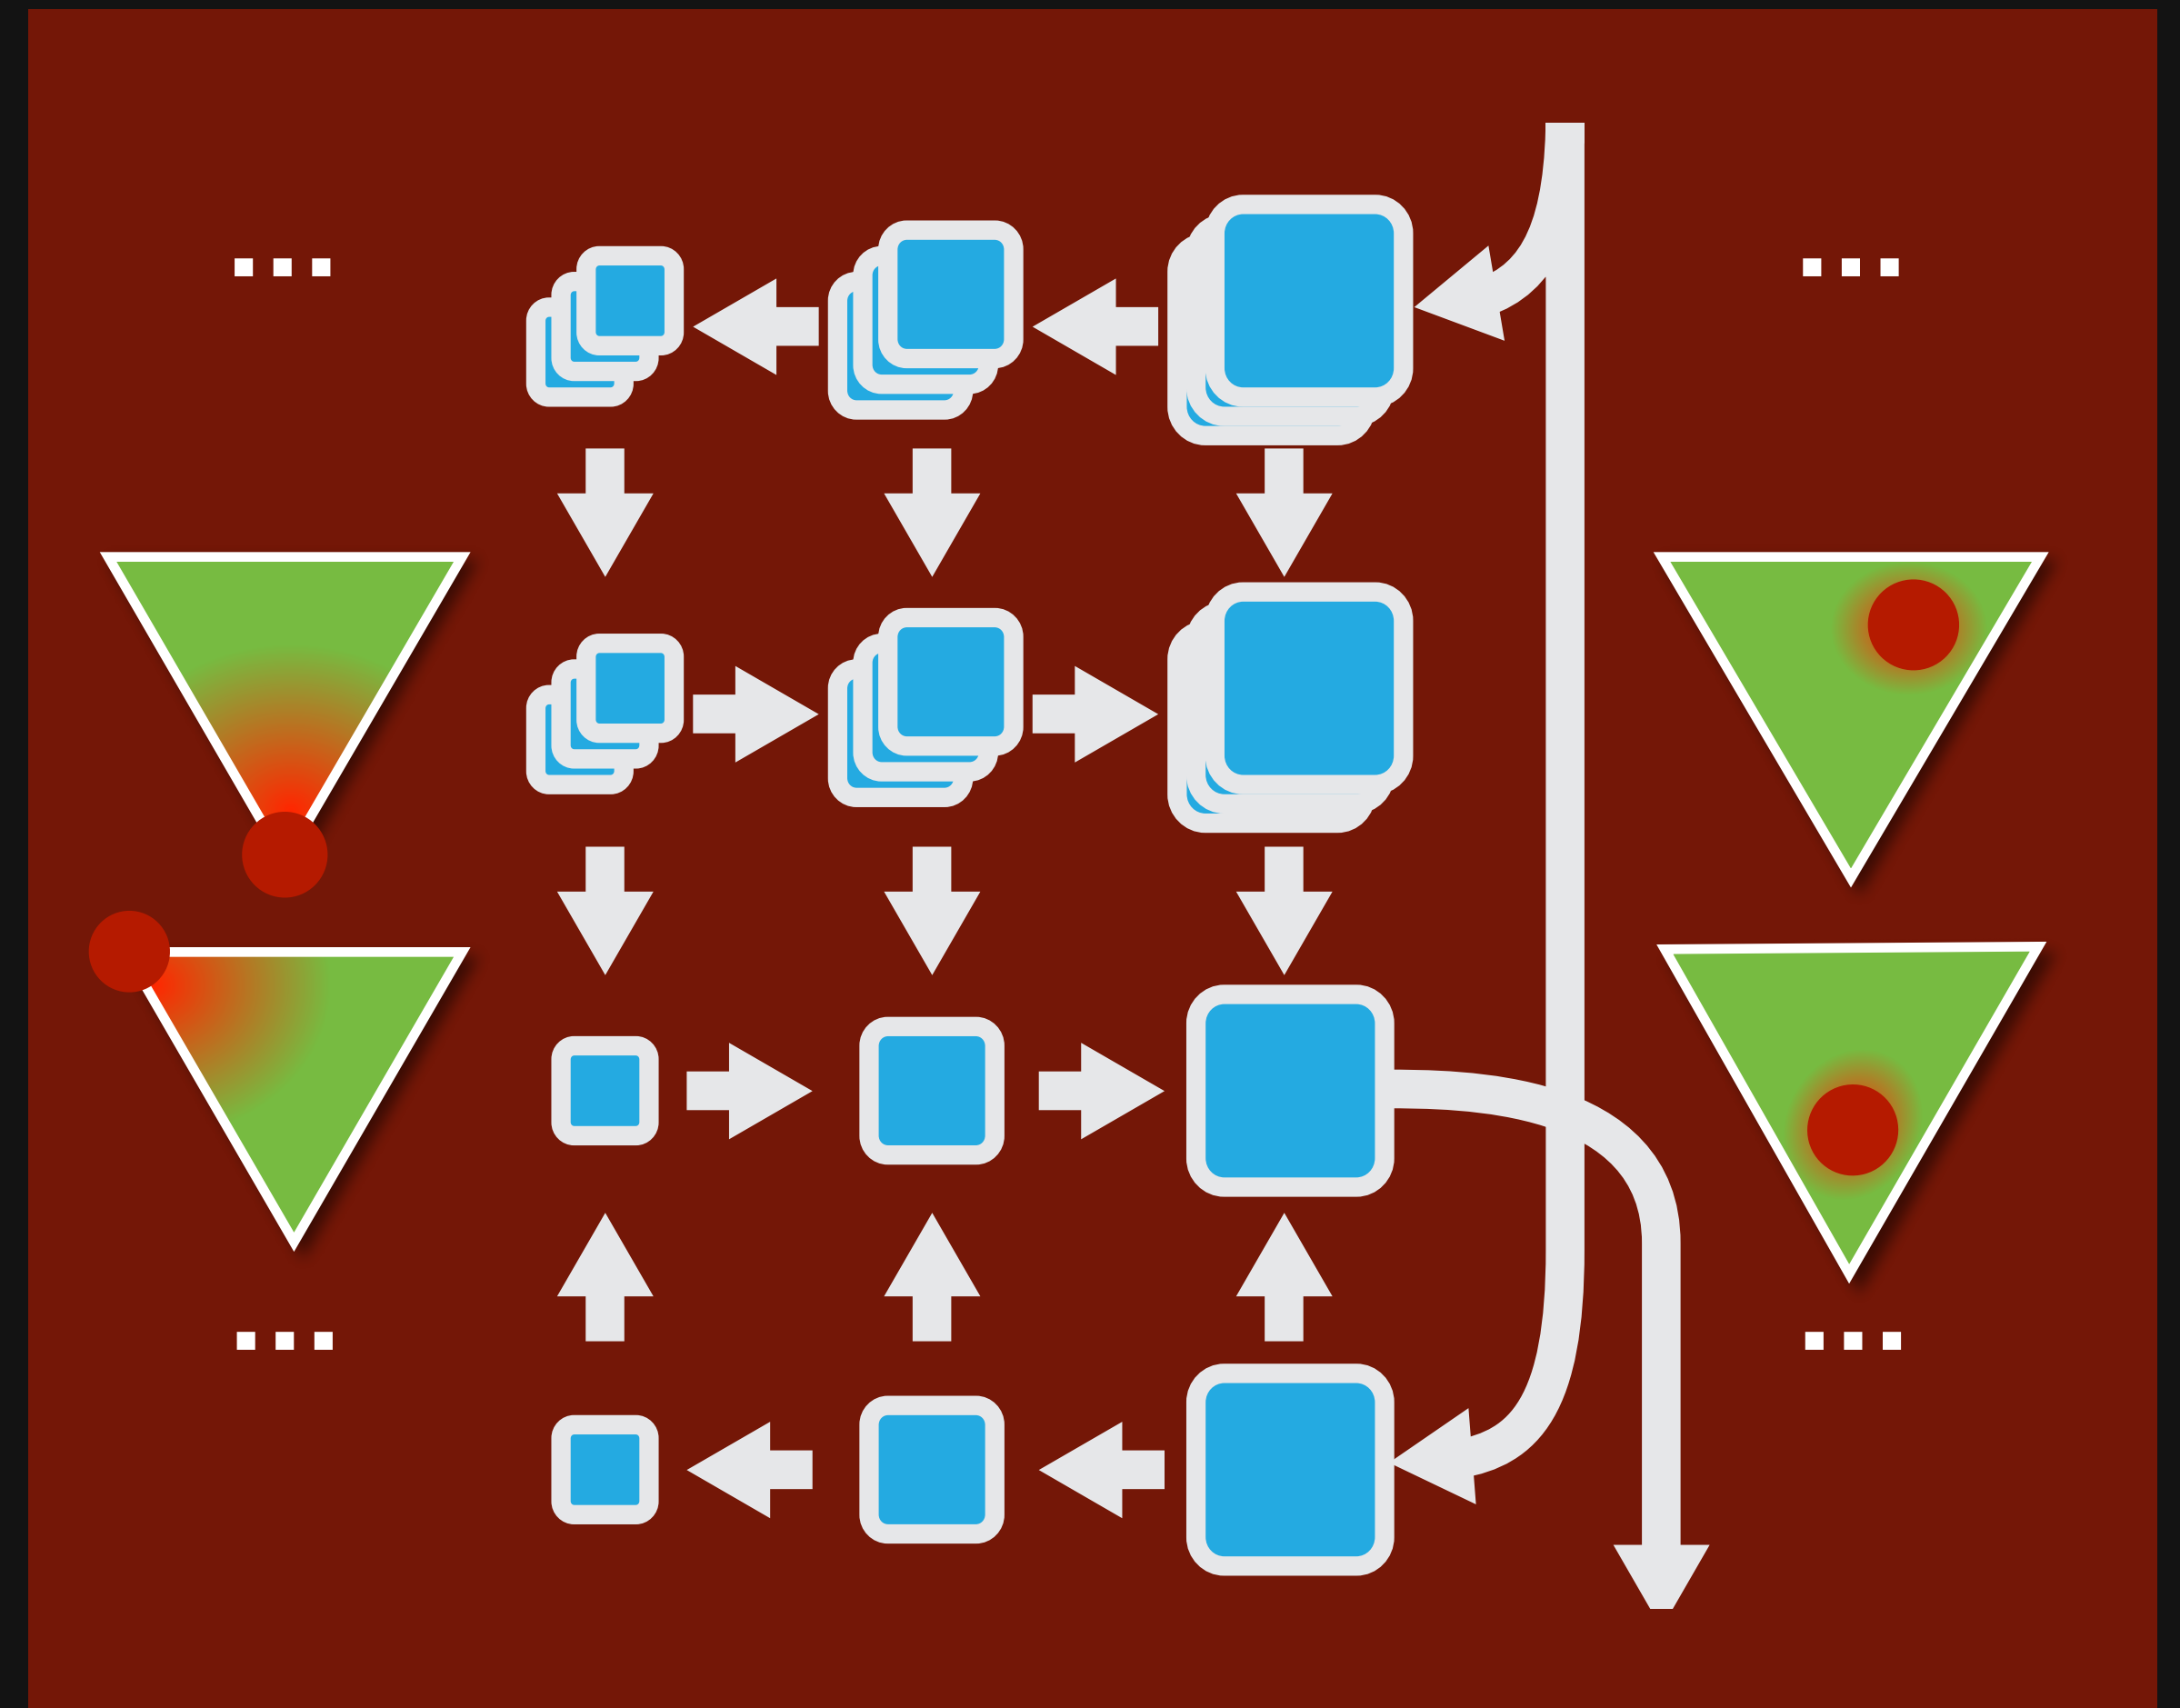
*1 intern-day*

**20x faster (vs. reference)**

**2x faster (vs. Adobe)**

**GPU: 90x faster (vs. reference)**

**9x faster (vs. Adobe)**



<b>x86</b>	Speedup	Factor shorter
Blur	1.2 ×	18 ×
Bilateral Grid	4.4 ×	4 ×
Camera pipeline	3.4 ×	2 ×
“Healing brush”	1.7 ×	7 ×
Local Laplacian	1.7 ×	5 ×

<b>GPU</b>	Speedup	Factor shorter
Bilateral Grid	2.3 ×	11 ×
“Healing brush”	5.9* ×	7* ×
Local Laplacian	9* ×	7* ×

<b>ARM</b>	Speedup	Factor shorter
Camera pipeline	1.1 ×	3 ×

<b>x86</b>	Speedup	Factor shorter
Blur	1.2 ×	18 ×
Bilateral Grid	4.4 ×	4 ×
Camera pipeline	3.4 ×	2 ×
“Healing brush”	1.7 ×	7 ×
Local Laplacian	1.7 ×	5 ×
Gaussian Blur	1.5 ×	5 ×
FFT (vs. FFTW)	1.5 ×	10s
BLAS (vs. Eigen)	1 ×	100s

<b>GPU</b>	Speedup	Factor shorter
Bilateral Grid	2.3 ×	11 ×
“Healing brush”	5.9* ×	7* ×
Local Laplacian	9* ×	7* ×

<b>ARM</b>	Speedup	Factor shorter
Camera pipeline	1.1 ×	3 ×

# Current status

open source at <http://halide-lang.org>

## Google

65 active developers

> 200 pipelines

10s of kLOC in production

**G Photos *auto-enhance***

Data center

Android

Chrome (PNaCl)

***n* secret/unannounced projects**



**HDR+**

Glass

Nexus devices



**Adobe**



**Movidius** 

>20 companies  
on Halide-Dev

# Today's agenda

the big ideas in Halide

Now: **writing & optimizing real code**

Hello world (brightness)

Gaussian blur - *3x OpenCV*

Simple enhancement pipeline - *6x OpenCV*

MATLAB integration

IIR filter

CNN layers

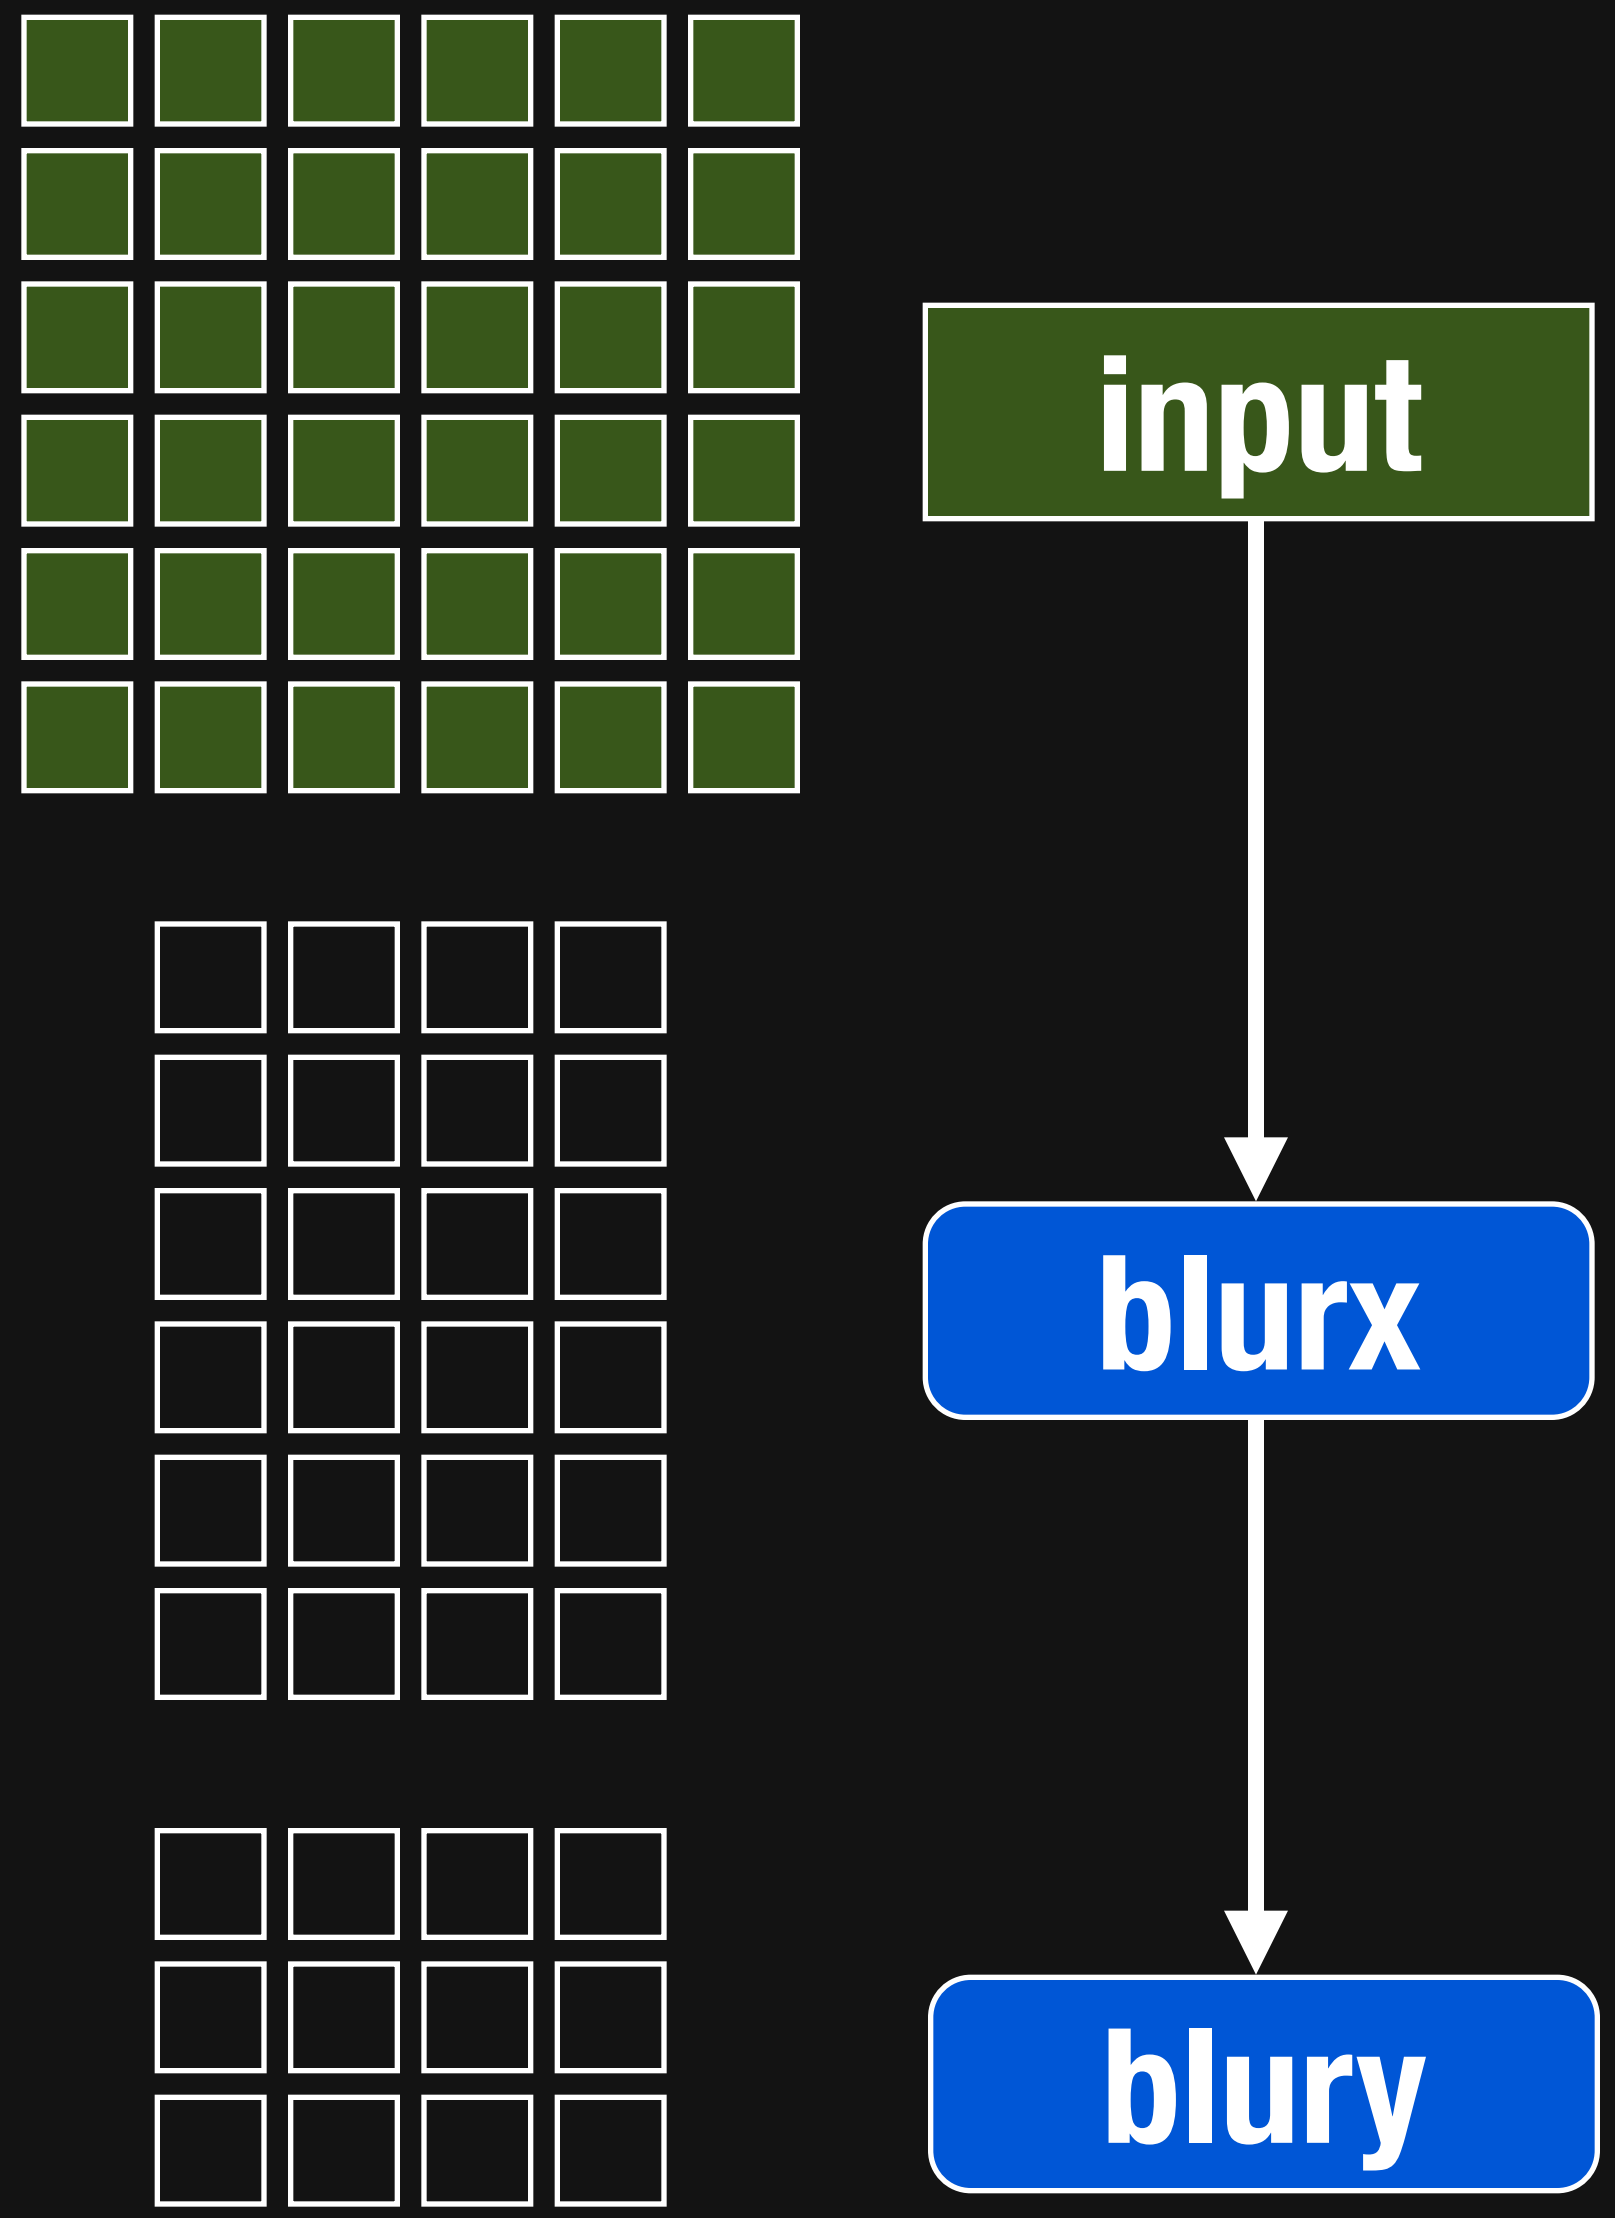
GPU scheduling

*break*

*break*

Finally: **real-time HOG on a phone**





# The **schedule** defines intra-stage order, inter-stage interleaving



# The **schedule** defines intra-stage order, inter-stage interleaving

For each stage:

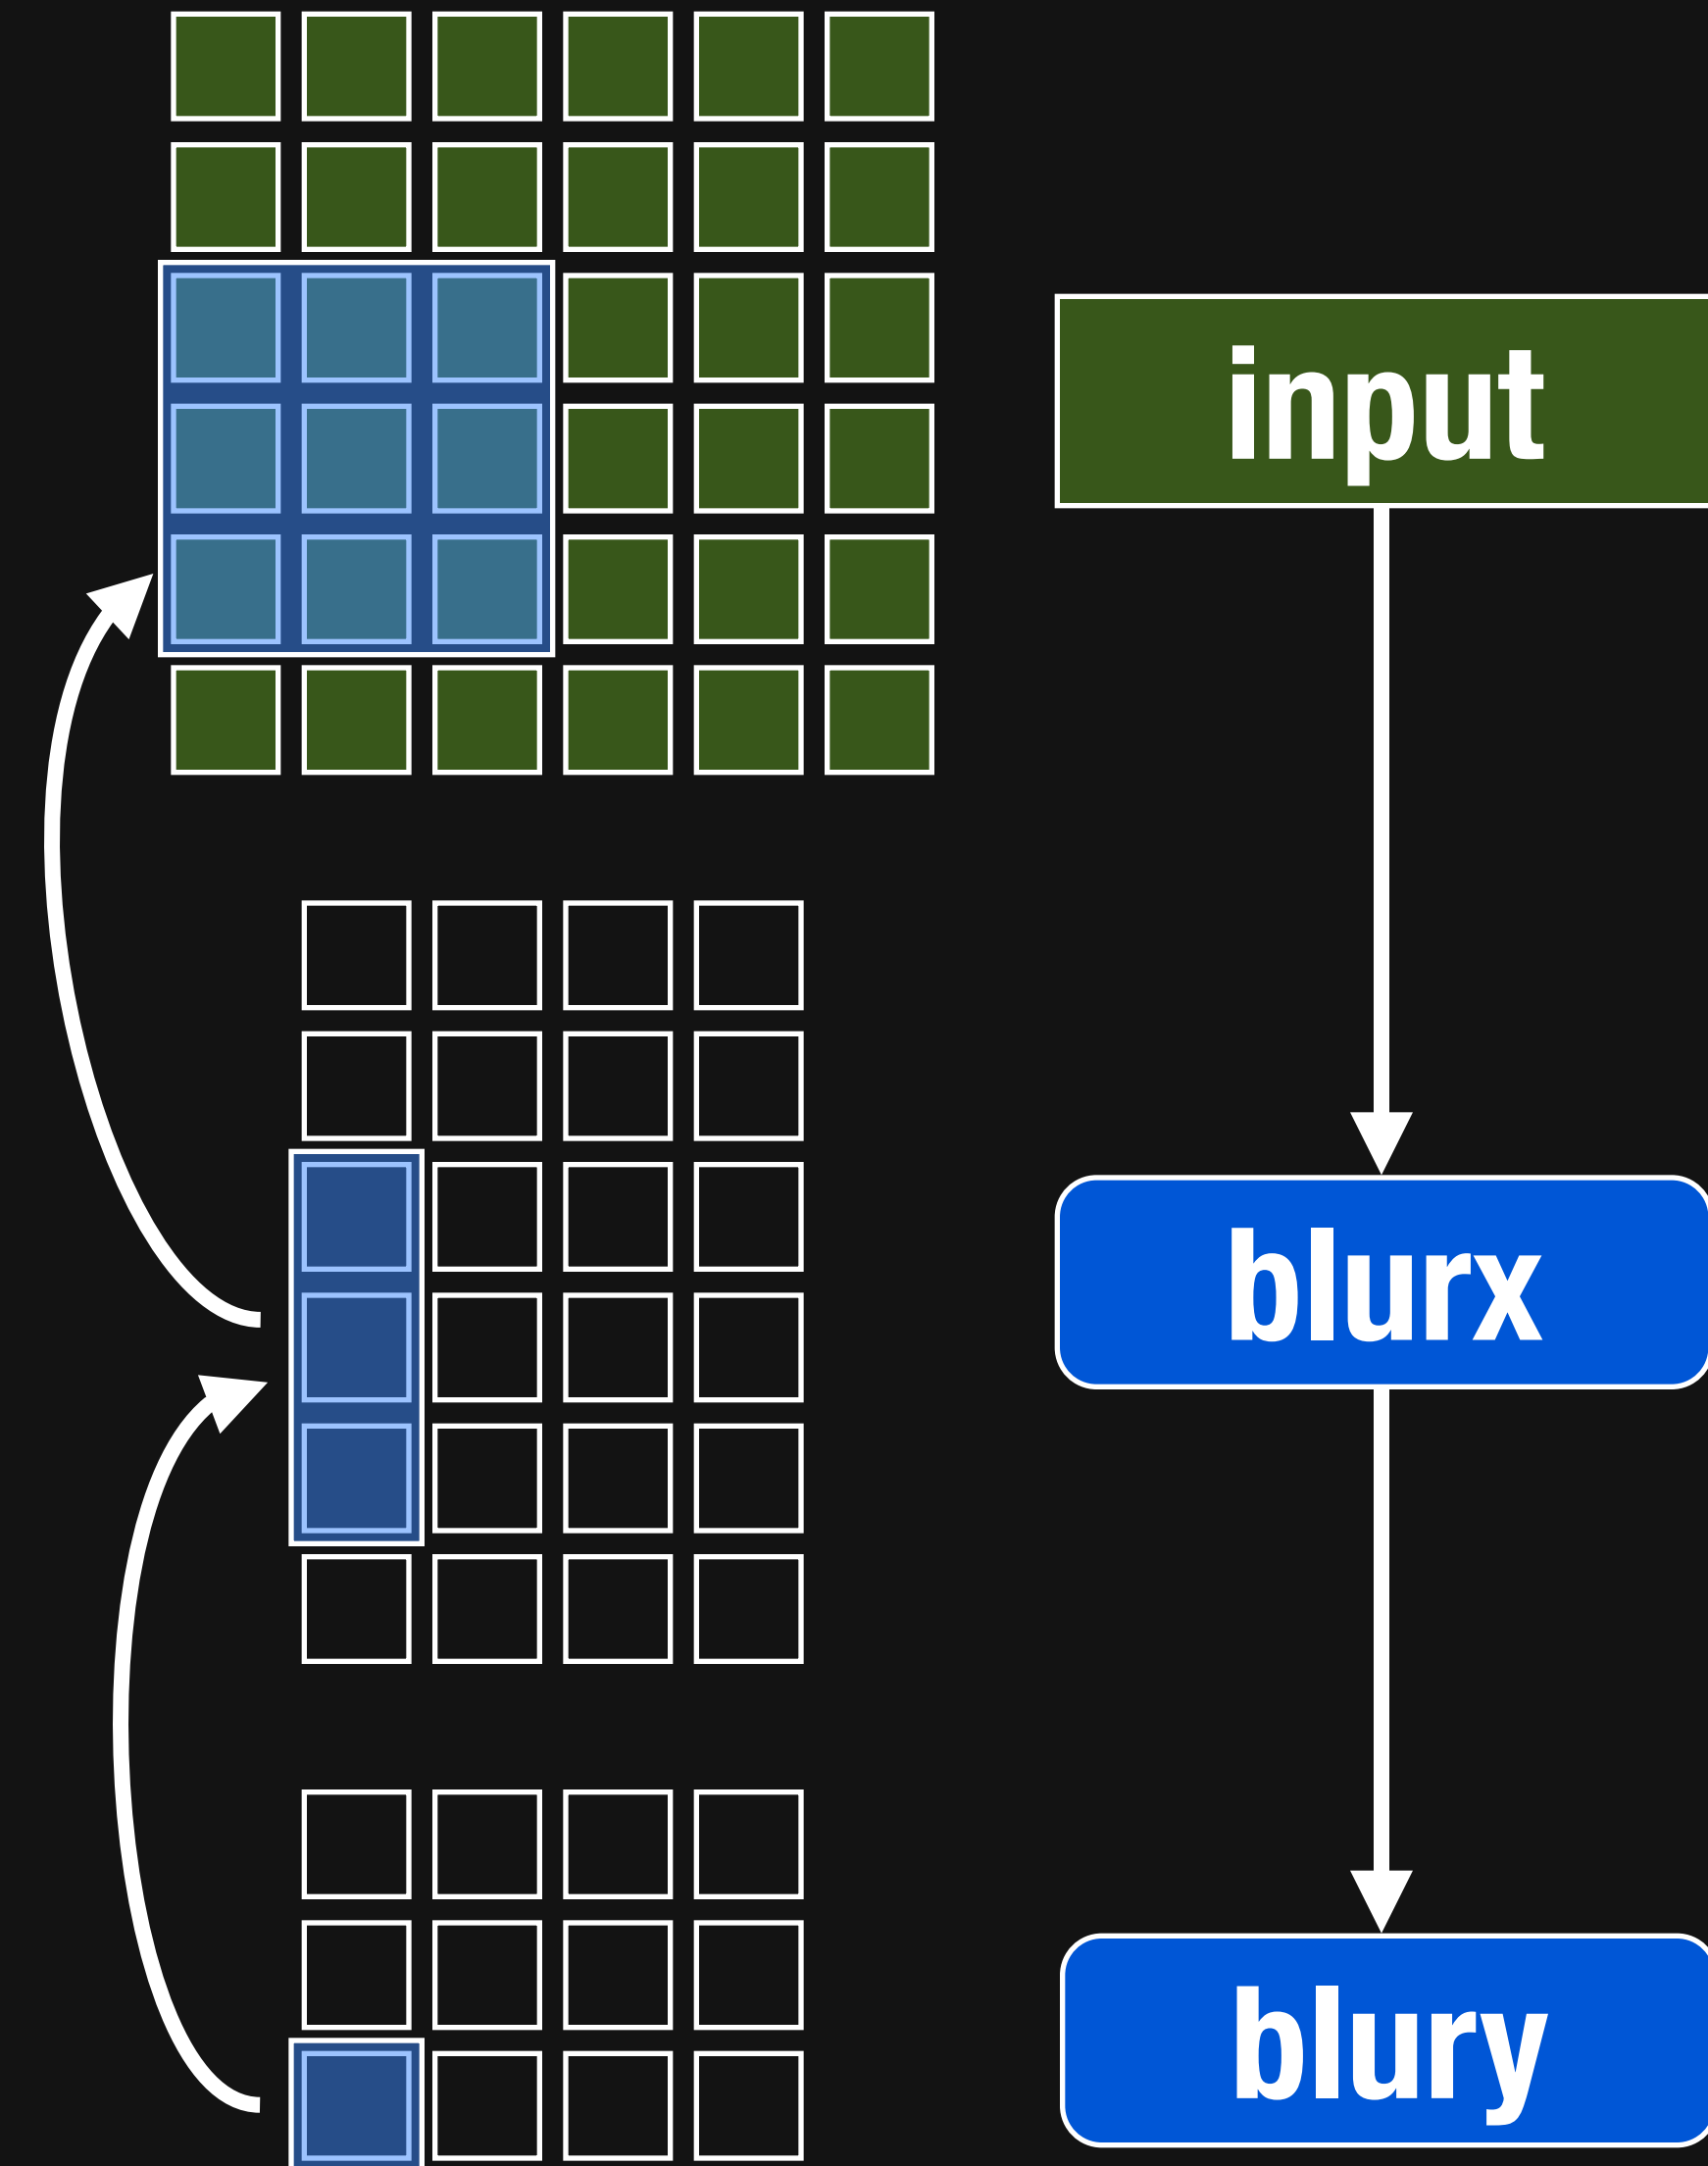
1) In what order should we compute its values?



# The **schedule** defines intra-stage order, inter-stage interleaving

**For each stage:**

- 1) In what order should we compute its values?
- 2) When should we compute its inputs?





**The schedule** defines order & parallelism within stages

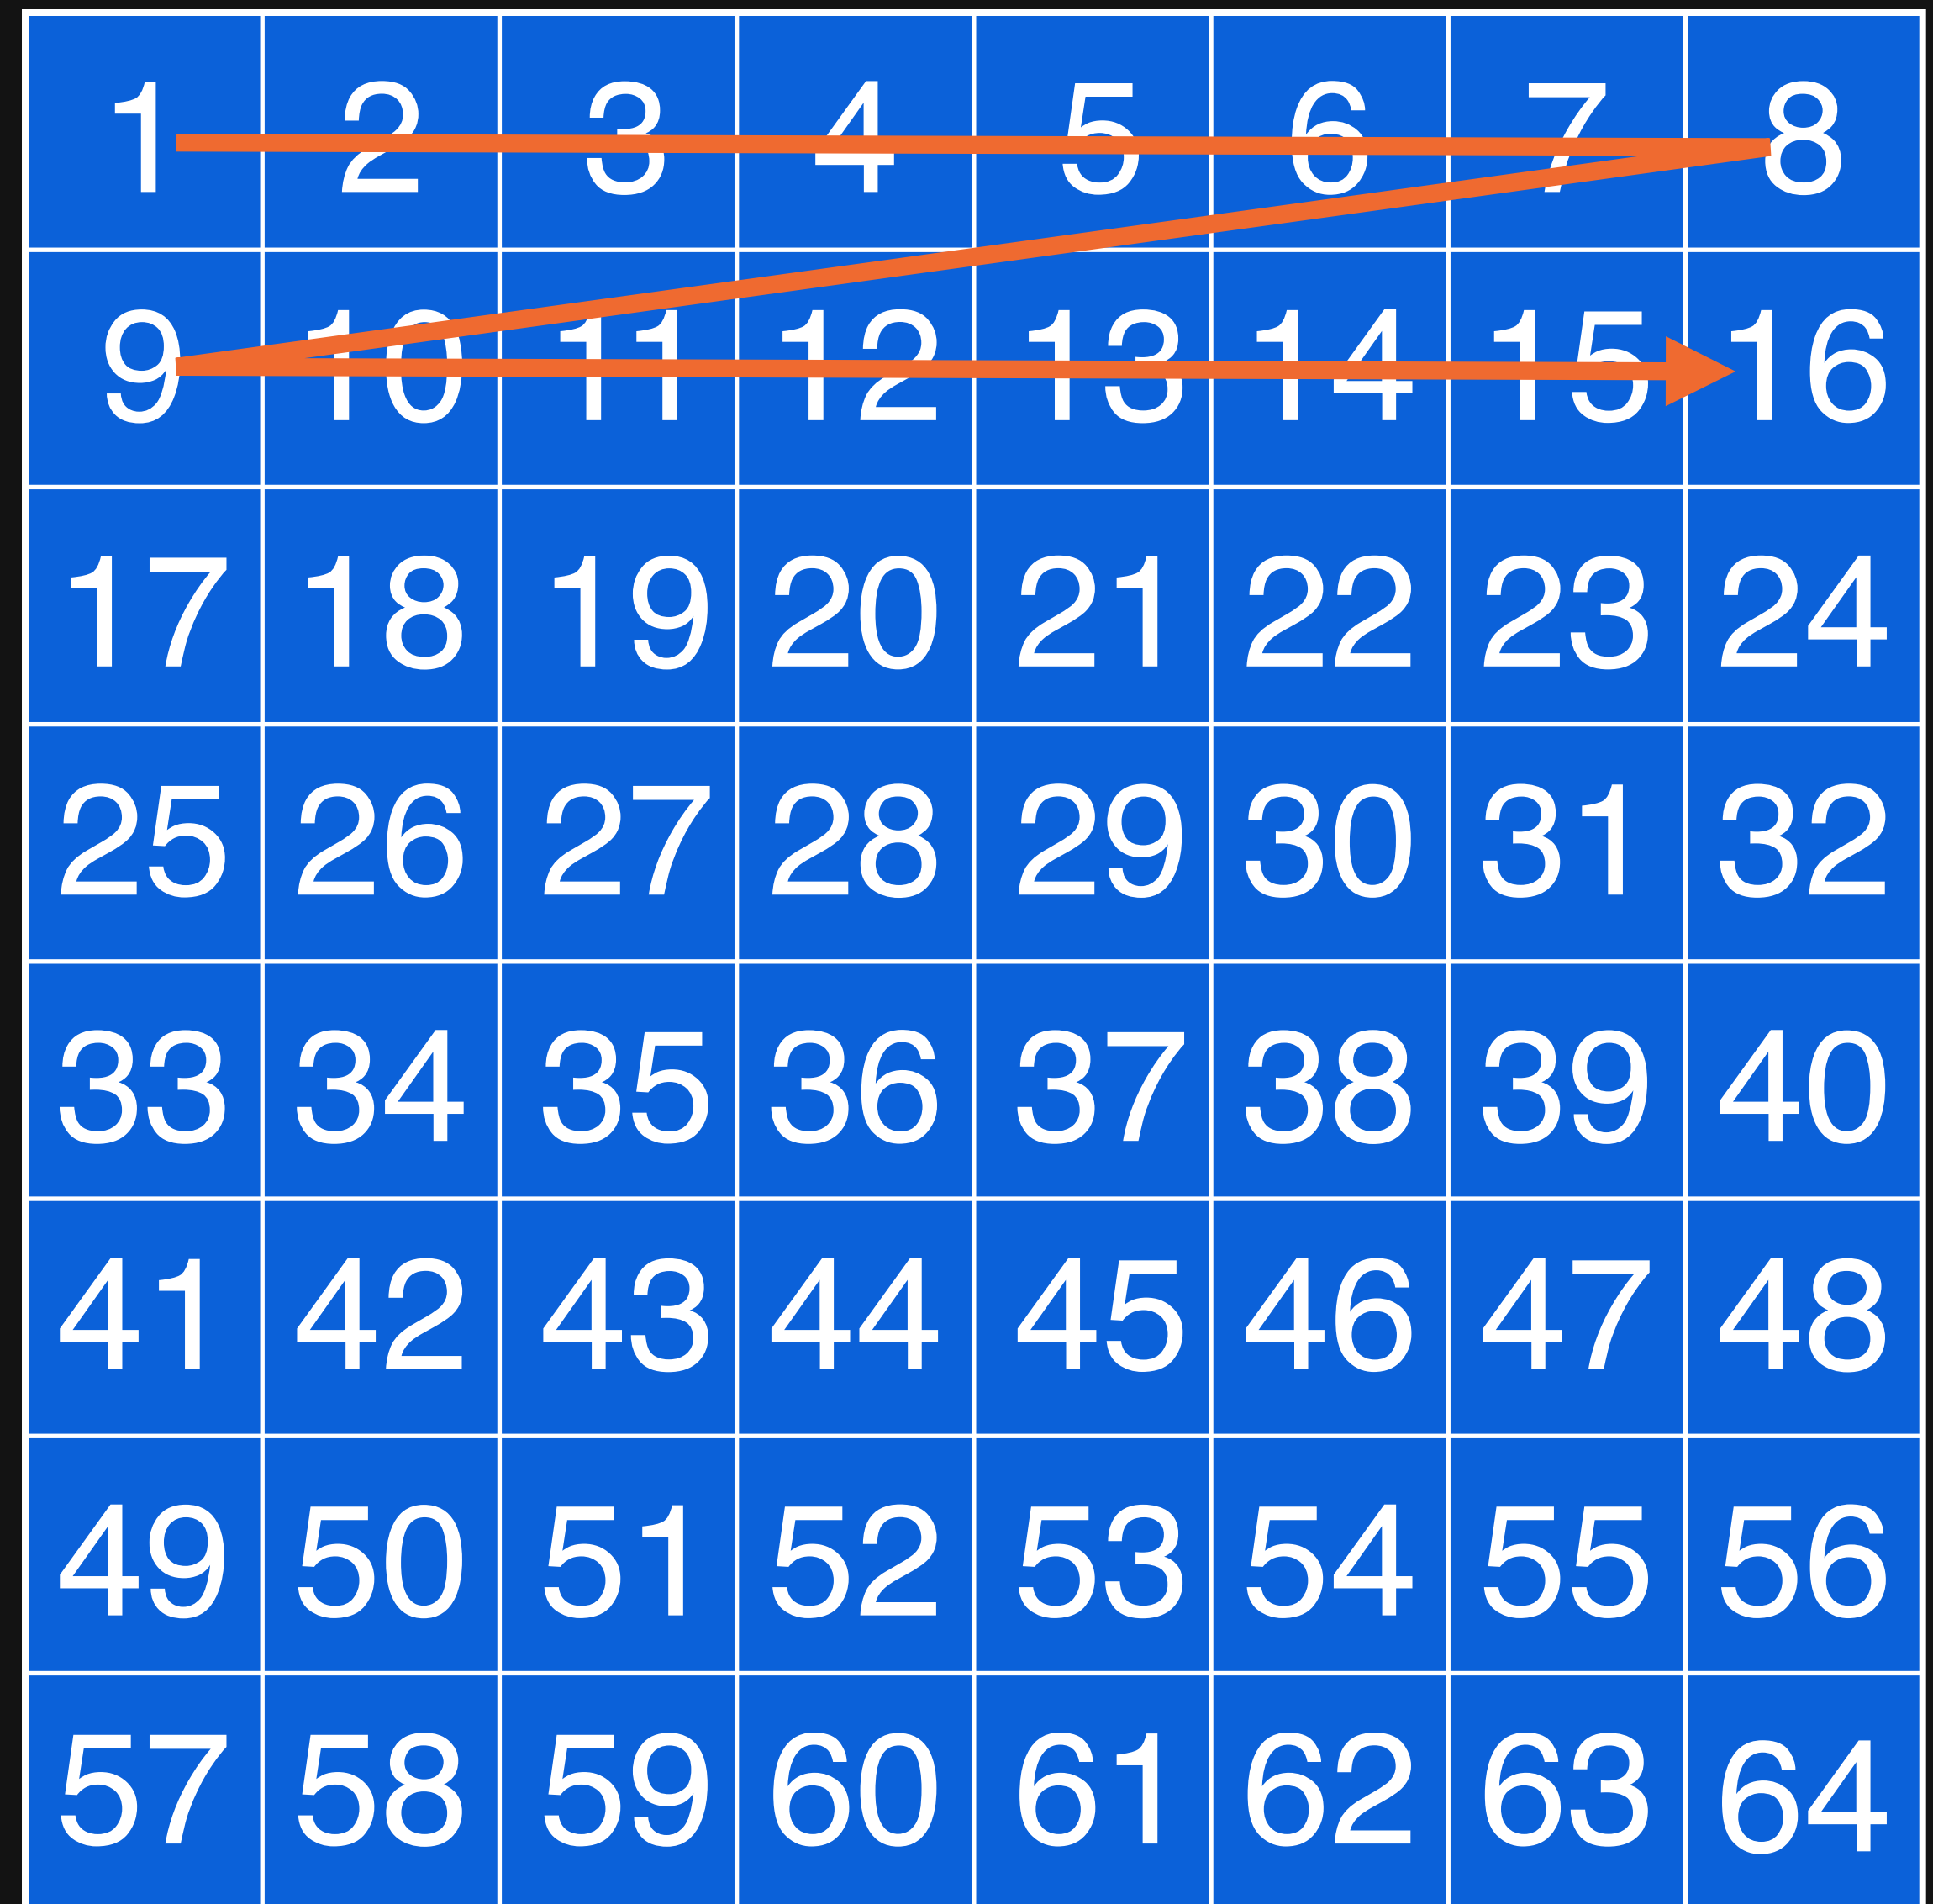
# The schedule defines order & parallelism within stages

**Serial y,  
Serial x**

1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22	23	24
25	26	27	28	29	30	31	32
33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48
49	50	51	52	53	54	55	56
57	58	59	60	61	62	63	64

# The schedule defines order & parallelism within stages

**Serial y,  
Serial x**



1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22	23	24
25	26	27	28	29	30	31	32
33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48
49	50	51	52	53	54	55	56
57	58	59	60	61	62	63	64



# The schedule defines order & parallelism within stages

**Serial y,  
Vectorize x by 4**

	1			2	
	3			4	
	5			6	
	7			8	
	9			10	
	11			12	
	13			14	
	15			16	

# The schedule defines order & parallelism within stages

**Serial y,  
Vectorize x by 4**

	1			2	
	3			4	
	5			6	
	7			8	
	9			10	
	11			12	
	13			14	
	15			16	







# The schedule defines order & parallelism within stages

**Split x by 2,  
Split y by 2.**

1	2	5	6	9	10	13	14
3	4	7	8	11	12	15	16
17	18	21	22	25	26	29	30
19	20	23	24	27	28	31	32
33	34	37	38	41	42	45	46
35	36	39	40	43	44	47	48
49	50	53	54	57	58	61	62
51	52	55	56	59	60	63	64

# The schedule defines order & parallelism within stages

**Split  $x$  by 2,**  
**Split  $y$  by 2.**  
**Serial  $y_{\text{outer}},$**   
**Serial  $x_{\text{outer}},$**   
**Serial  $y_{\text{inner}},$**   
**Serial  $x_{\text{inner}}$**

1	2	5	6	9	10	13	14
3	4	7	8	11	12	15	16
17	18	21	22	25	26	29	30
19	20	23	24	27	28	31	32
33	34	37	38	41	42	45	46
35	36	39	40	43	44	47	48
49	50	53	54	57	58	61	62
51	52	55	56	59	60	63	64

Domain order defines a **loop nest** for each function

Domain order defines a **loop nest** for each function

brighten(x, y, c) = ...



Domain order defines a **loop nest** for each function

```
brighten(x, y, c) = ...  
    for c:  
        for y:  
            for x:  
                brighten(...) = ...
```

**Default:**

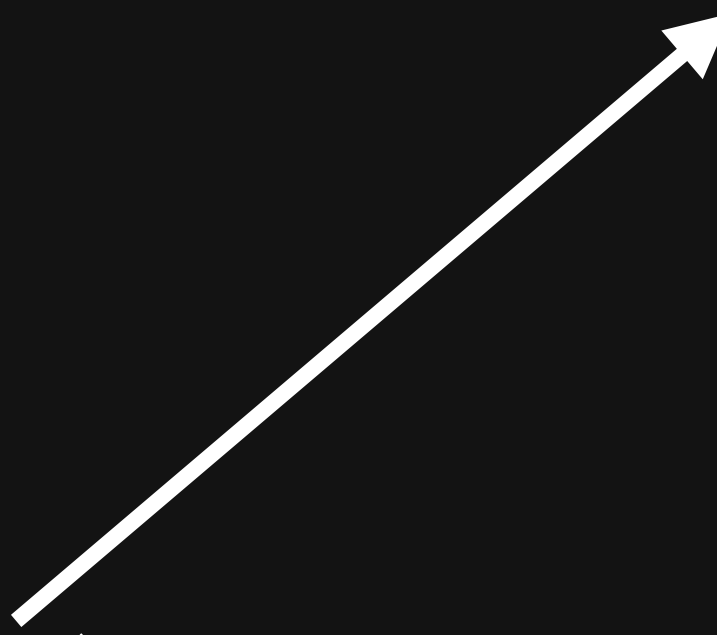
Serial c,

Serial y,

Serial x

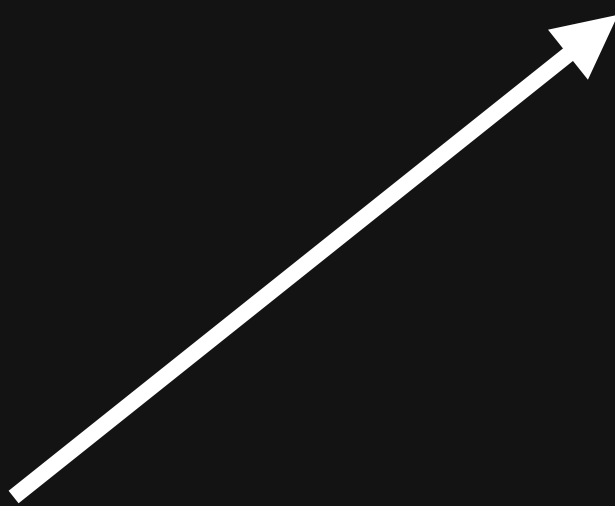
# Parallel marks a loop to be multithreaded

```
brighten(x, y, c) = ...  
  
brighten.parallel(y)  
    for c:  
        parallel y:  
            for x:  
                brighten(...) = ...
```




# Parallel marks a loop to be multithreaded

```
brighten(x, y, c) = ...  
    for c:  
        parallel y:  
            for x:  
                vectorized x.v in [0,7]:  
                    brighten(...) = ...  
brighten.parallel(y)  
    .vectorize(x, 8)
```



# Parallel marks a loop to be multithreaded

```
brighten(x, y, c) = ...  
    for c:  
        parallel y:  
            for x:  
                unrolled x.v in [0,3]:  
                    brighten(...) = ...  
brighten.parallel(y)  
    .unroll(x, 4)
```

A white arrow points from the `.unroll(x, 4)` method call in the bottom code block to the `unrolled x.v in [0,3]:` line in the top code block, illustrating how the unroll operation is implemented by unrolling the inner loop.



# Parallel marks a loop to be multithreaded


```
brighten(x, y, c) = ...  
    for c:  
        for yo:  
            for yi in [0,63]:  
                for x:  
                    brighten(...) = ...  
brighten.split(y, yo, yi, 64)
```

The diagram illustrates the relationship between a function call and its implementation. The function call `brighten.split(y, yo, yi, 64)` is shown at the bottom left. Two white arrows originate from this call: one points to the `for yo:` loop header, and the other points to the `for yi in [0,63]:` loop header. The nested loop structure is shown to the right, with the `for` keywords highlighted in blue. The function call `brighten(x, y, c) = ...` is shown at the top left, with the `for c:` loop header also highlighted in blue.

# Parallel marks a loop to be multithreaded

```
brighten(x, y, c) = ...  
  
brighten.split(y, y0, yi, 64)  
    .reorder(c, y0)
```

`for` y<sub>0</sub>:  
 `for` c:  
 `for` y<sub>i</sub> in [0,63]:  
 `for` x:  
 brighten(...) = ...



# Today's agenda

the big ideas in Halide

writing & optimizing real code

Now: Hello world (brightness)

Gaussian blur - *3x OpenCV*

Simple enhancement pipeline - *6x OpenCV*

MATLAB integration

IIR filter

CNN layers

GPU scheduling

*break*

*break*

Finally: **real-time HOG on a phone**

# Today's agenda

the big ideas in Halide

writing & optimizing real code

Hello world (brightness)

Now: Gaussian blur

Simple enhancement pipeline - *6x OpenCV*

MATLAB integration

IIR filter

CNN layers

GPU scheduling

*break*

*break*





Finally: **real-time HOG on a phone**



# The **schedule** defines intra-stage order, inter-stage interleaving

For each stage:

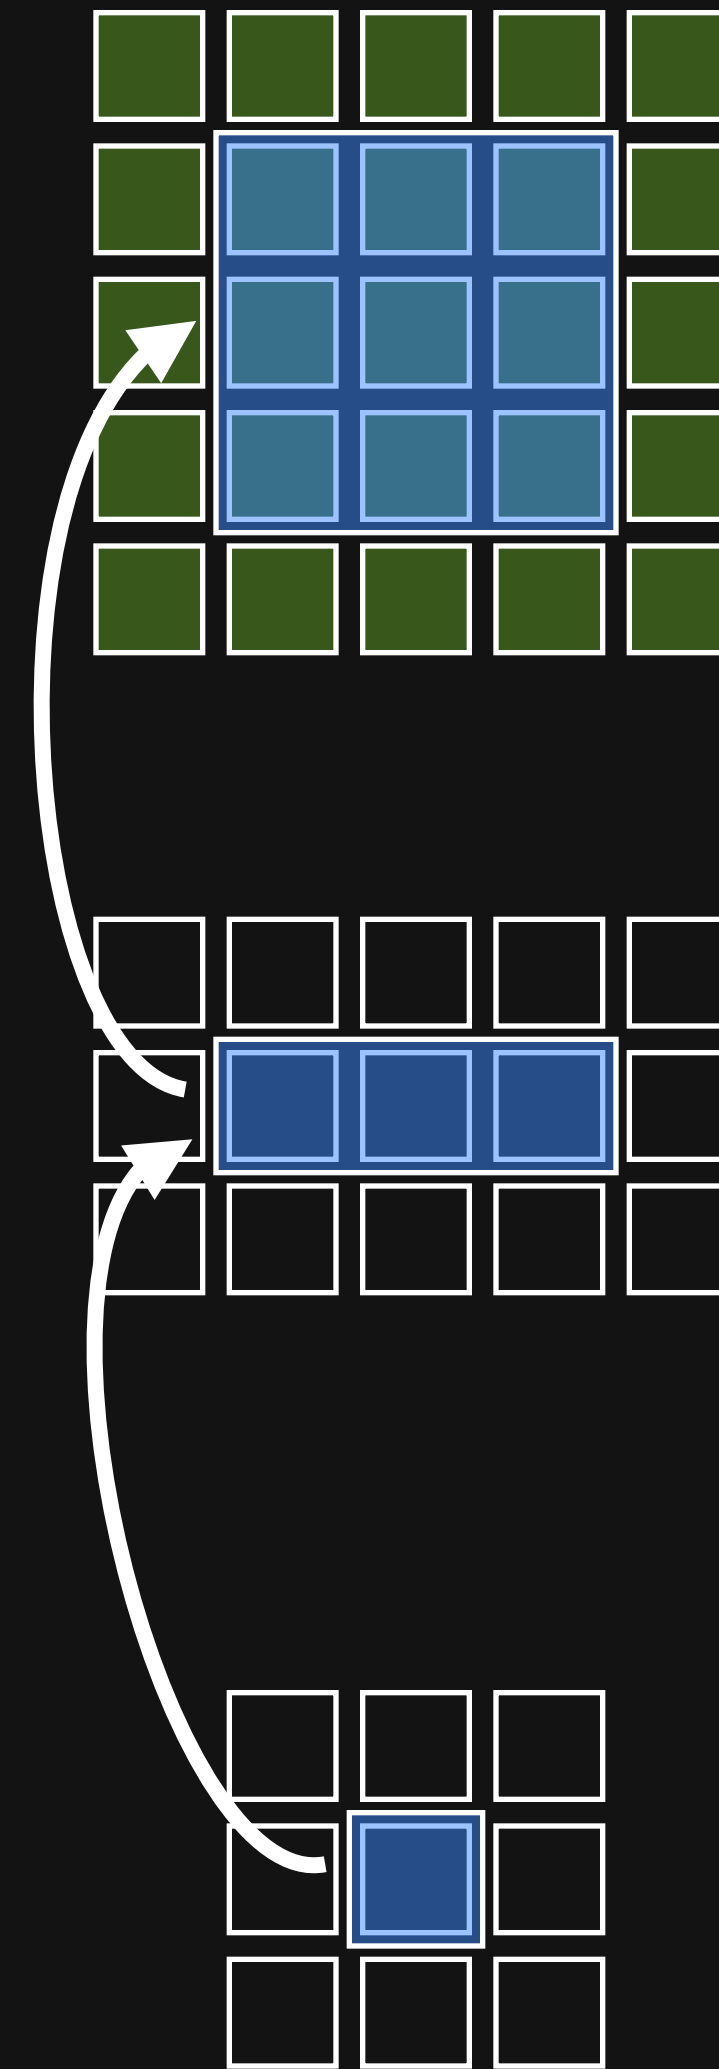
- 1) In what order should we compute its values?
- 2) When should we compute its inputs?




# The **schedule** defines intra-stage order, inter-stage interleaving

For each stage:

- 1) In what order should we compute its values?
- 2) When should we compute its inputs?



# Organizing the algorithm as a **data-parallel pipeline** & **loops**




# Organizing the algorithm as a data-parallel pipeline & loops

input

blurry

blurx





# Organizing the algorithm as a **data-parallel pipeline** & loops

input

blurry

blurx

```
compute blurry:  
  for ...:  
    blurry(...) = ...
```

```
compute blurx:  
  for ...:  
    blurx(...) = ...
```

# Inline maximizes locality, but also recomputes values

input

blurry

blurx


```
compute blurx:
```

```
  for c:
```

```
    for y:
```

```
      for x:
```

```
        blurx(...) =
```



# Inline maximizes locality, but also recomputes values




```
compute blurx:
```

```
  for c:
```

```
    for y:
```

```
      for x:
```

```
        blurx(...) =
```



# Compute root minimizes **recompute**, but also **locality**



```
compute blur:
```

```
  for c:
```

```
    for y:
```

```
      for x:
```

```
        blur(...) = ...
```

```
compute blurx:
```


```
  for c:
```

```
    for y:
```

```
      for x:
```

```
        blurx(...) = ...
```

# Compute root minimizes **recompute**, but also **locality**



```
compute blur:
```

```
  for c:
```

```
    for y:
```

```
      for x:
```

```
        blur(...) = ...
```

```
compute blurx:
```

```
  for c:
```


```
    for y:
```

```
      for x:
```

```
        blurx(...) = ...
```



**Compute at blurx.y interleaves scanlines for better locality**



```
compute blurx:
```

```
  for c:
```

```
    for y:
```

```
      compute blury:
```

```
        for x:
```

```
          blury(...) = ...
```

```
      for x:
```

```
        blurx(...) = ...
```

**Compute at blurx.y interleaves scanlines for better locality**



```
compute blurx:
```

```
  for c:
```

```
    for y:
```

```
      compute blury:
```

```
        for x:
```

```
          blury(...) = ...
```

```
      for x:
```

```
        blurx(...) = ...
```

# Today's agenda

the big ideas in Halide

writing & optimizing real code

Hello world (brightness)

Now: Gaussian blur - *3x OpenCV*

Simple enhancement pipeline - *6x OpenCV*

MATLAB integration

IIR filter

CNN layers

GPU scheduling

*break*

*break*

Finally: **real-time HOG on a phone**